Convolutional Neural Network(CNN)

Các bài toán ứng dụng CNN

- Computer Vision

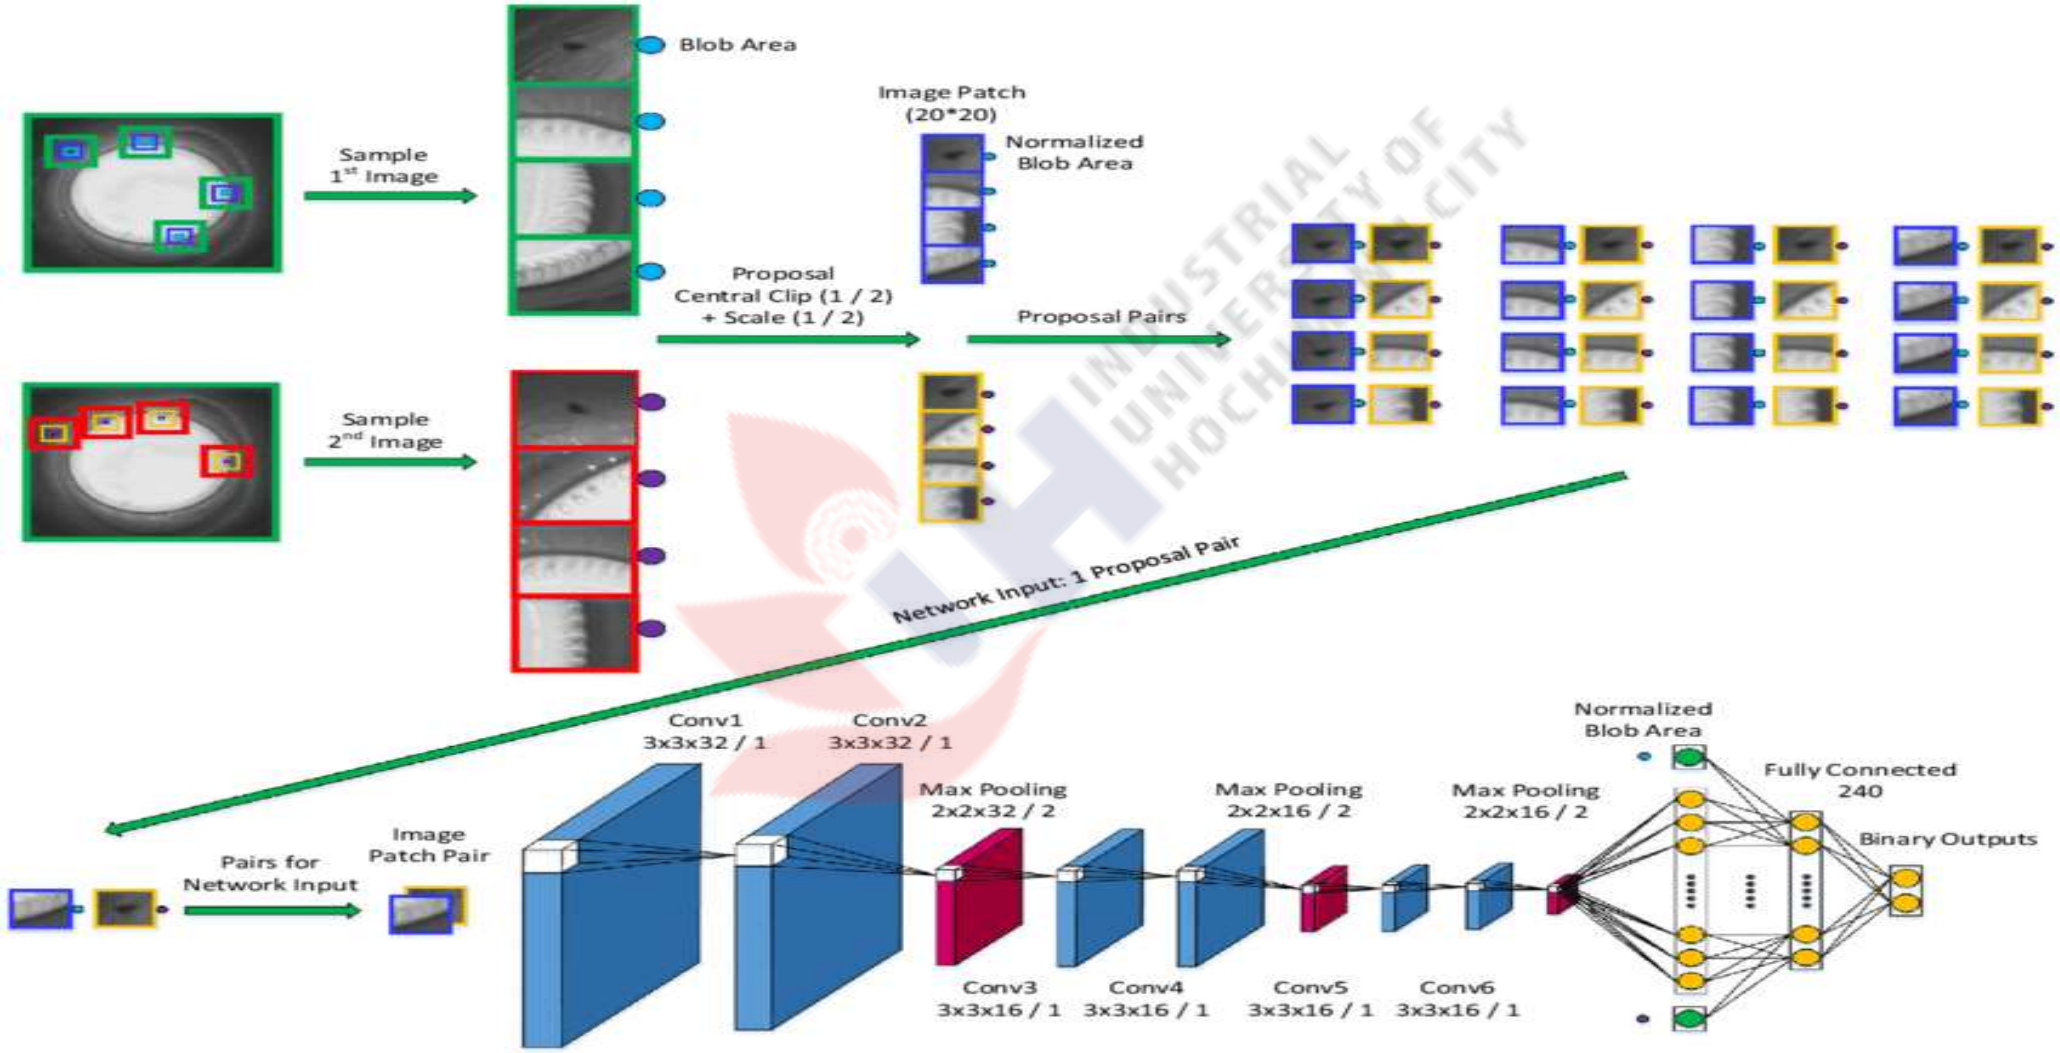

Các bài toán ứng dụng CNN

- NLP

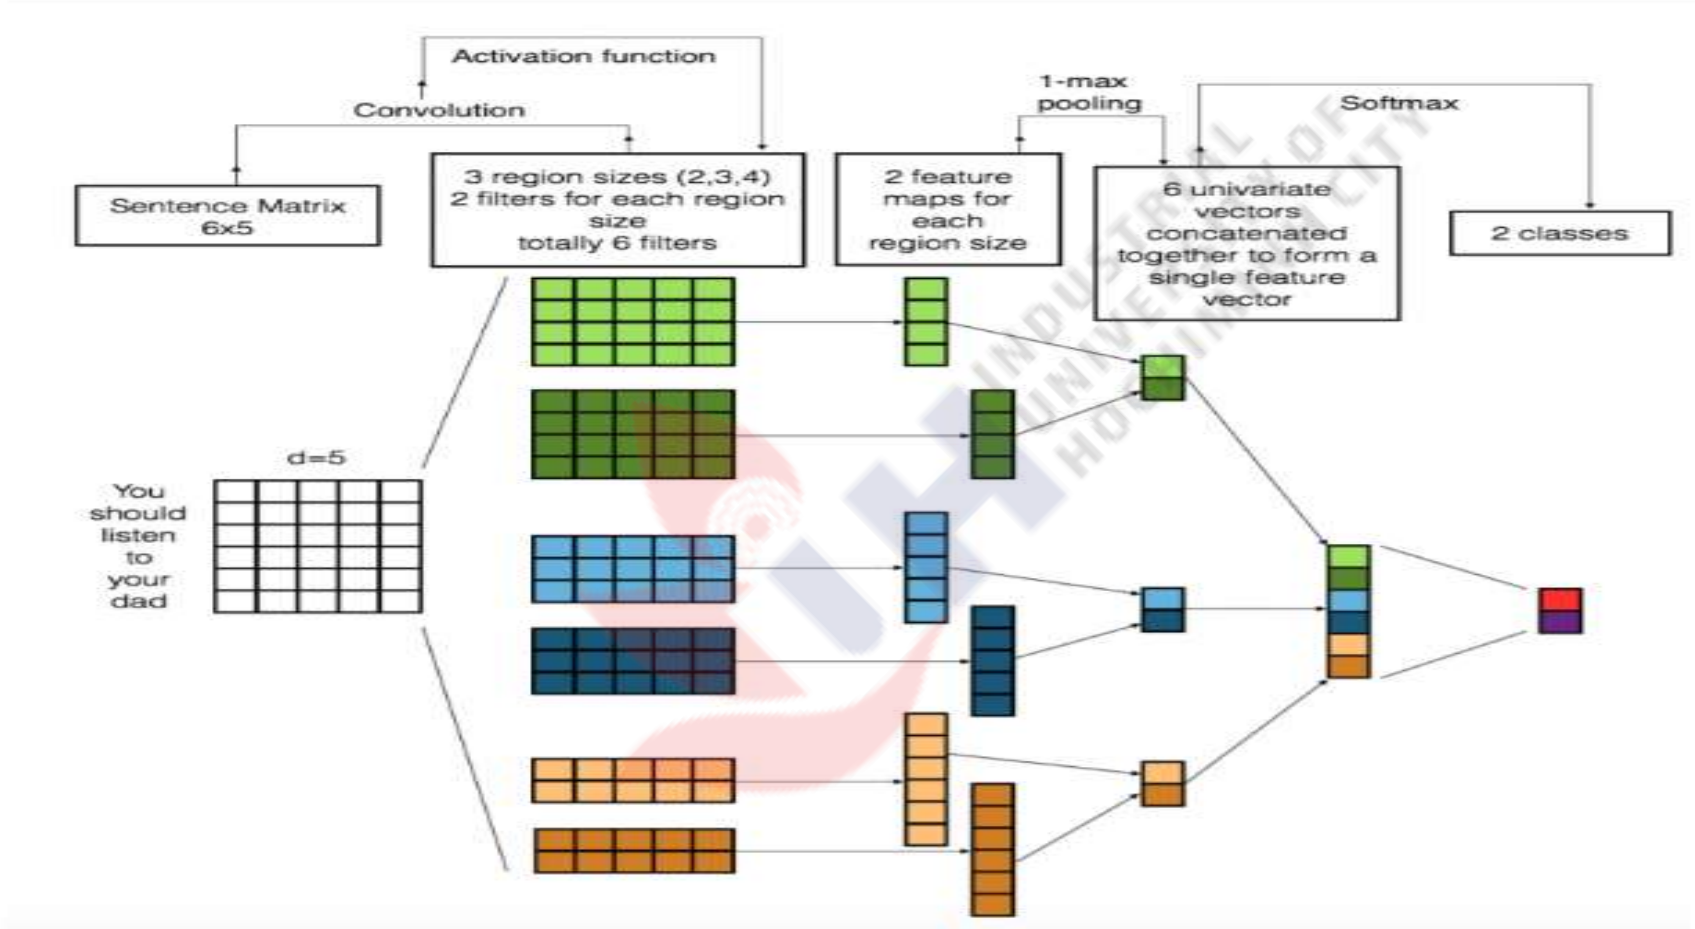

Computer Vision

- Feature extraction

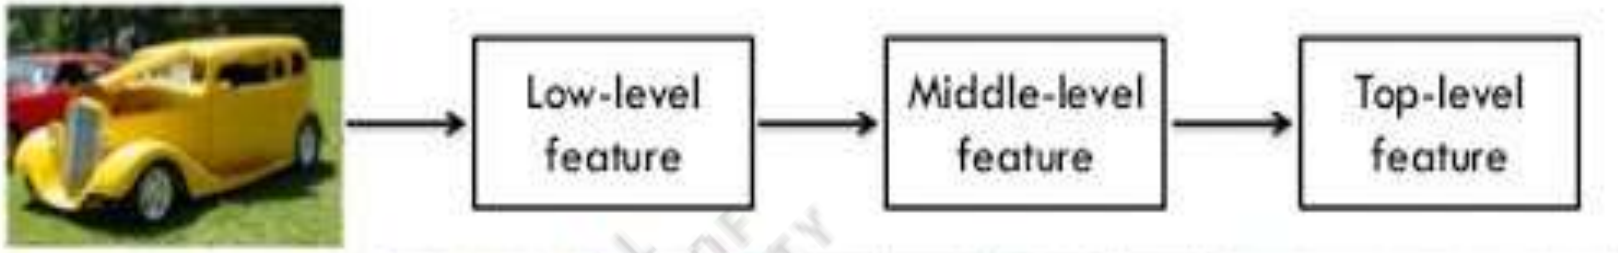

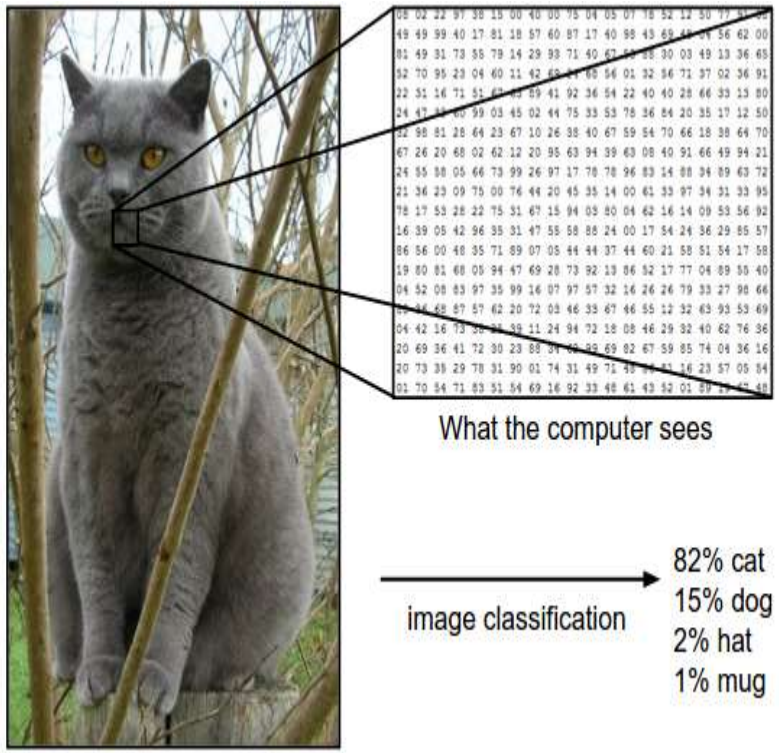

Feature extraction

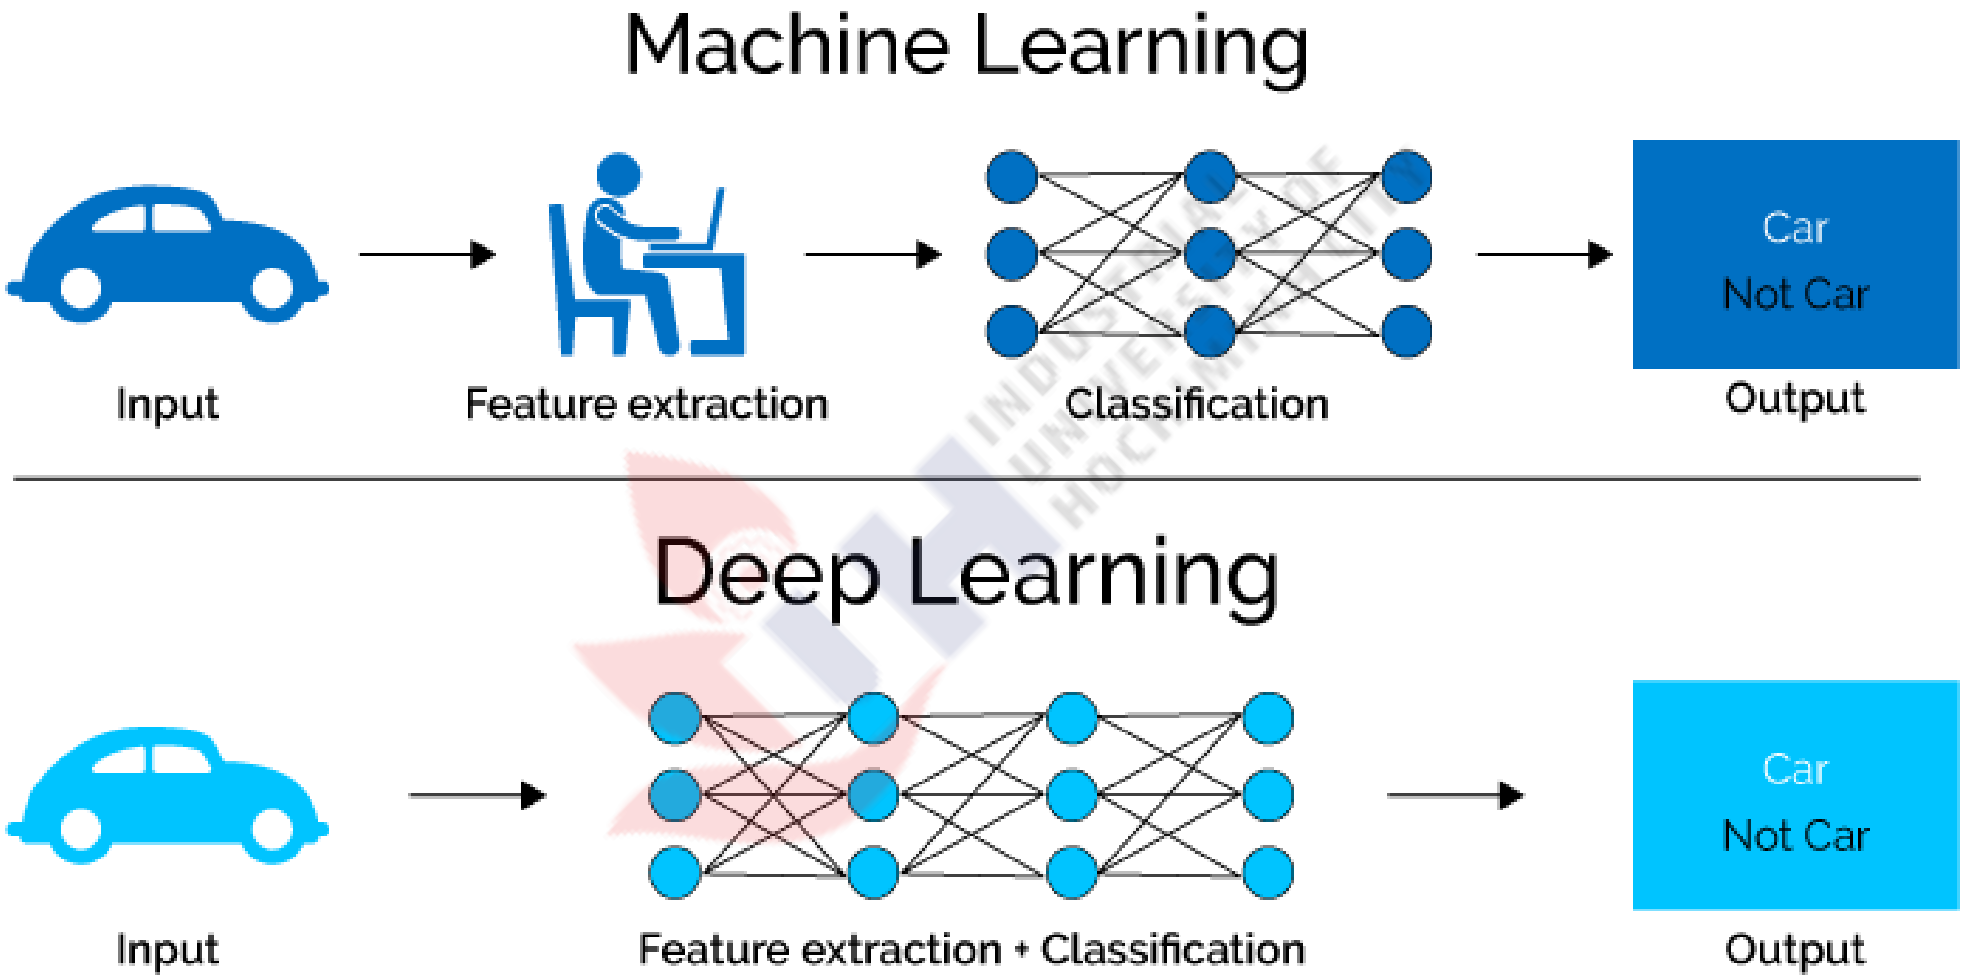

The Problem Space

Image: 480 x 480 x 3
Pixel : 0 to 255
Problem when 20.000 cats, 40.000 dogs, 50.000 birds???

What We See

What Computers See

Xf= 12288, số lượng nodes trong layer 1= 1000
=>Số lượng weight = 12288x1000=12288000
Số lượng bias= 12289000
Layer 2,…layer n???
⇒*Số lượng quá lớn, tính toán chậm, cần giải pháp tốt hơn!*
Áp dụng phép tính convolution vào các layer trong network giải quyết được vấn đề lượng lớn parameter mà vẫn lấy ra được các đặc trưng của ảnh.

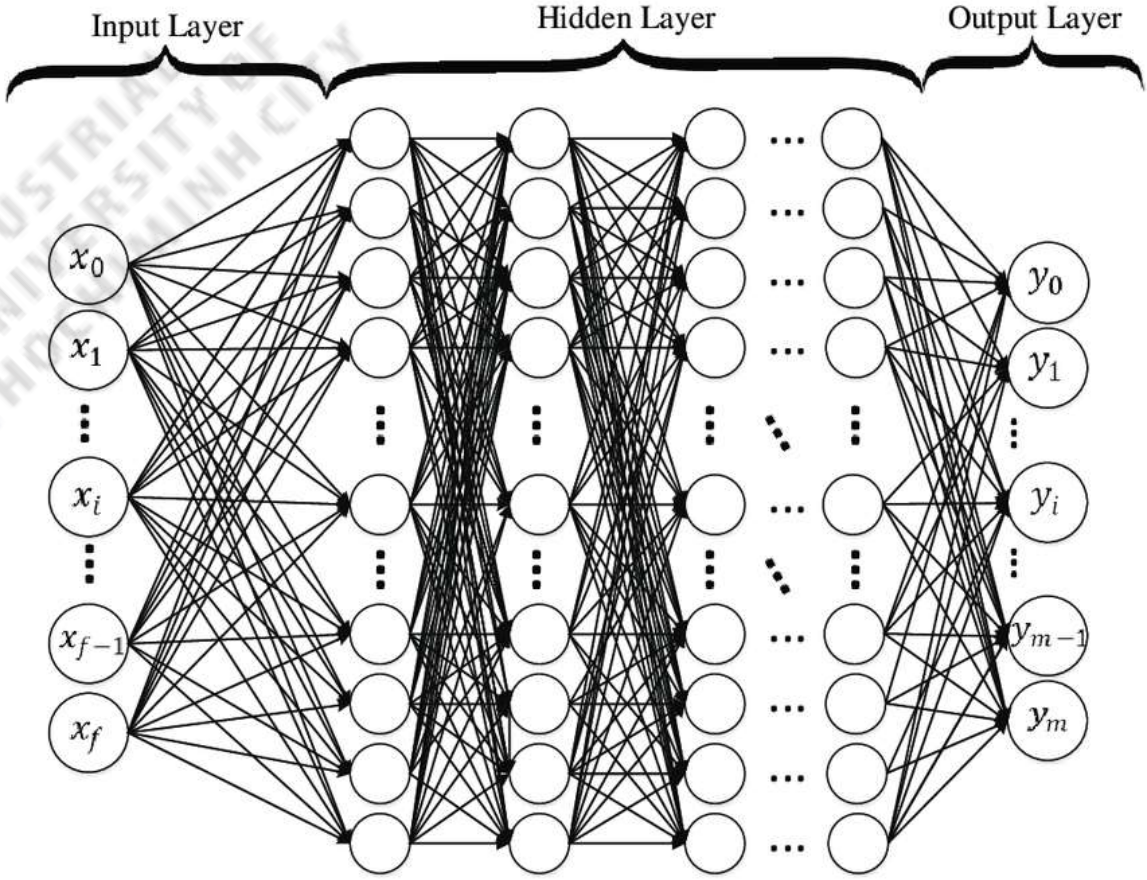

Convolutional Neural Network

Neural Networks receive an input (a single vector)
3D volumes of neurons: the layers of a ConvNet have neurons arranged in 3 dimensions: **width, height, depth**

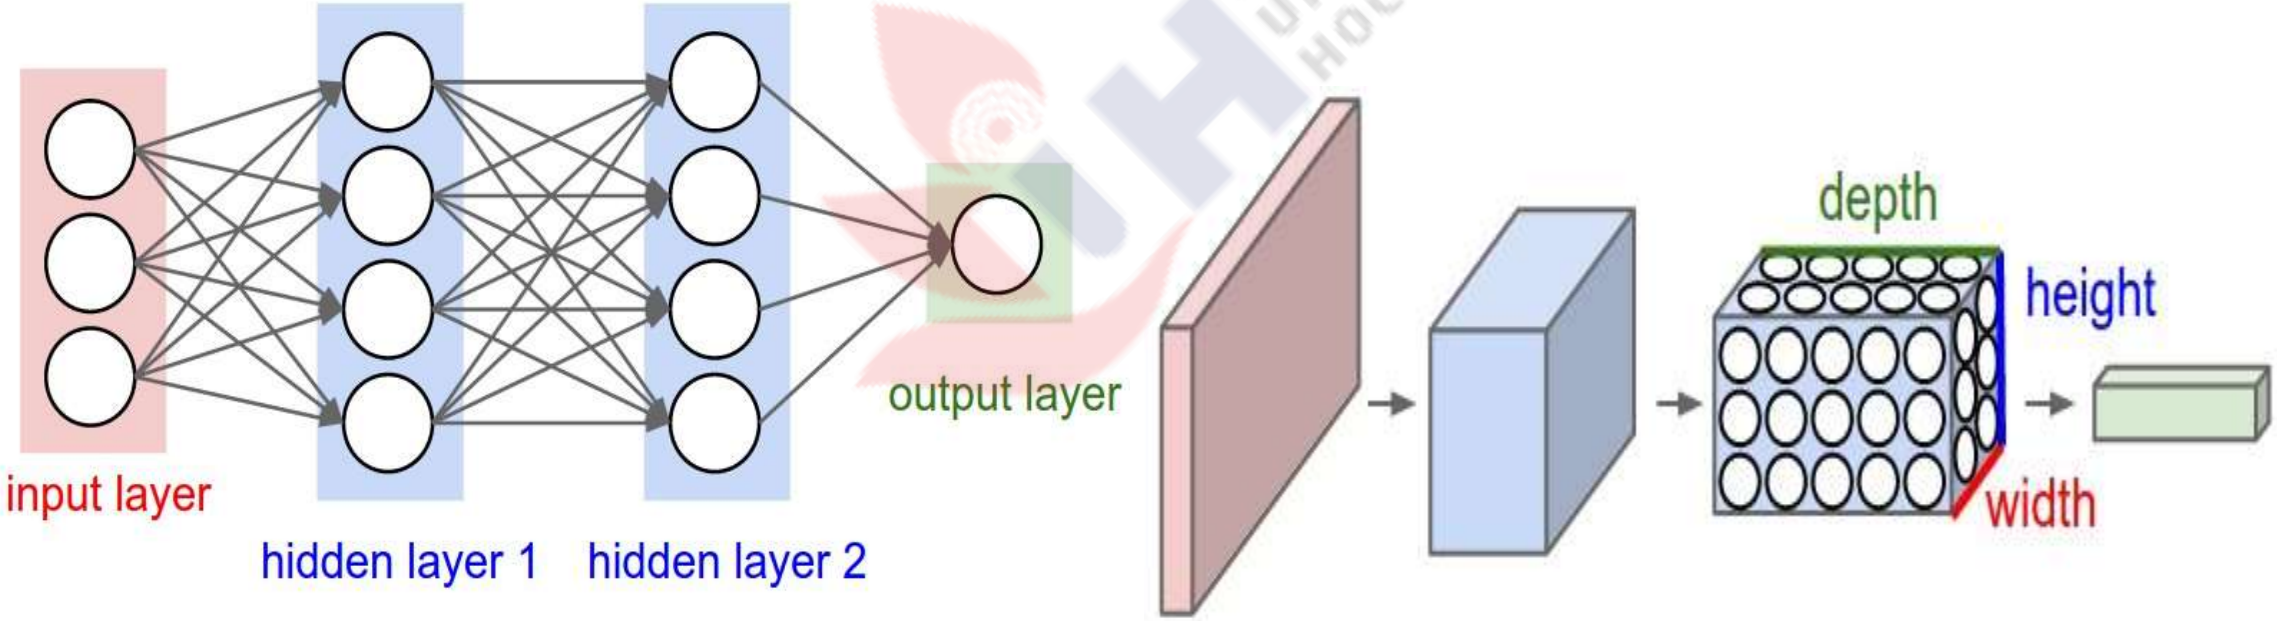

Convolutional Layer

A *"convolution"* is one of the building blocks of the Convolutional network.

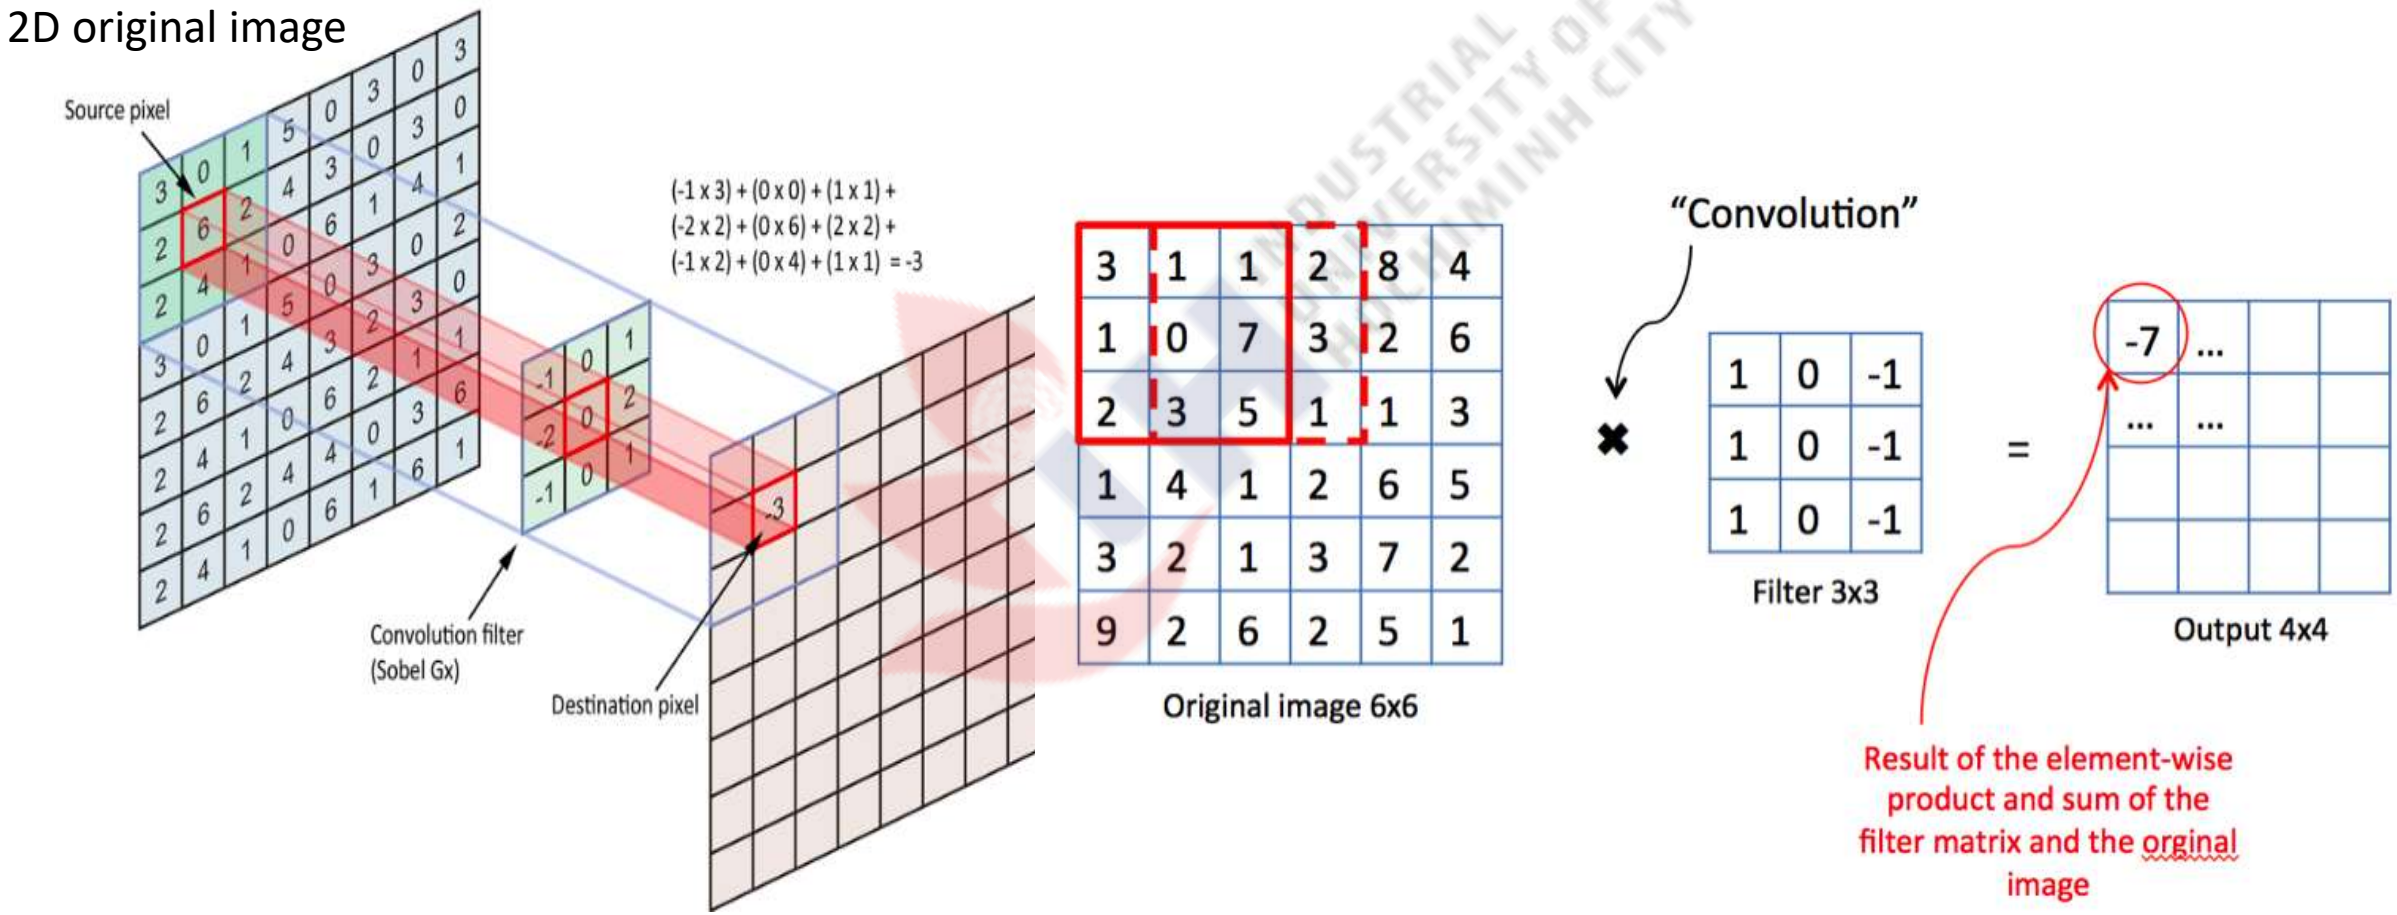

Step 1:

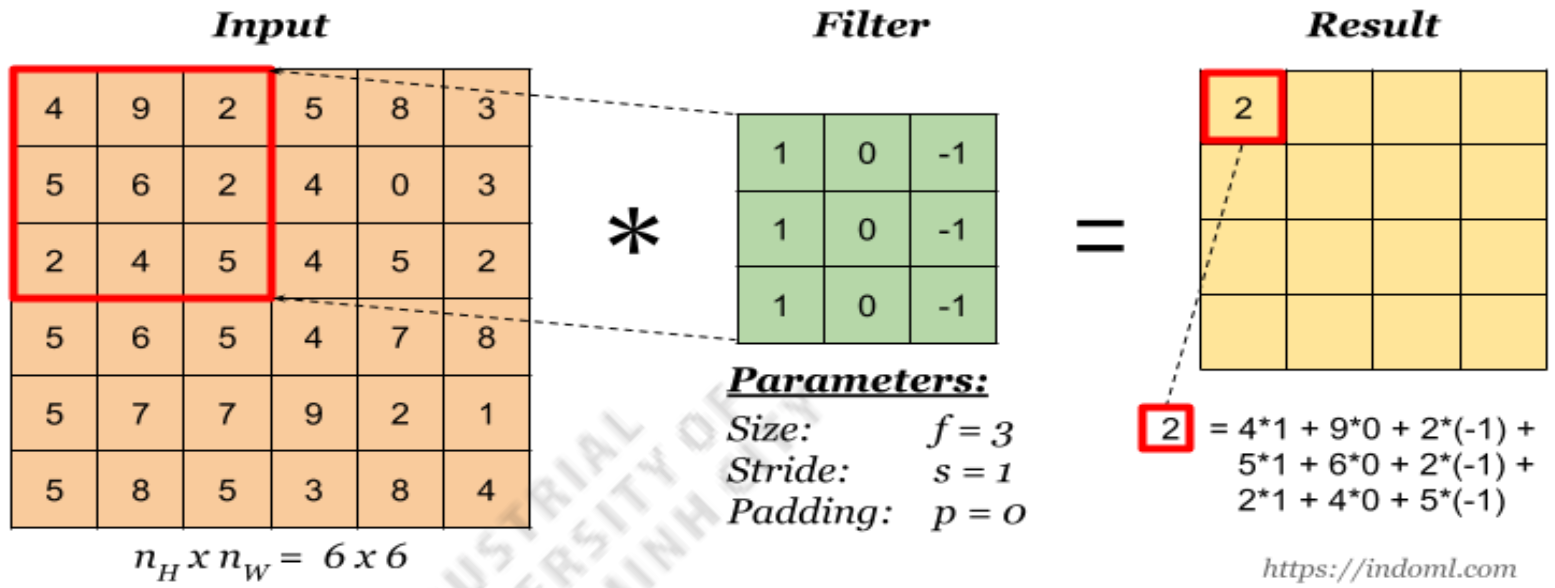
Input
Filter
Result
4 9 2 5 8 3
5 6 2 4 0 3
2 4 5 4 5 2
5 6 5 4 7 8
5 7 7 9 2 1
5 8 5 3 8 4
1 0 -1
1 0 -1
1 0 -1
2
Parameters:
Size: f = 3
Stride: s = 1
Padding: p = 0
2 = 4*1 + 9*0 + 2*(-1) +
5*1 + 6*0 + 2*(-1) +
2*1 + 4*0 + 5*(-1)
https://indoml.com

Step 2:

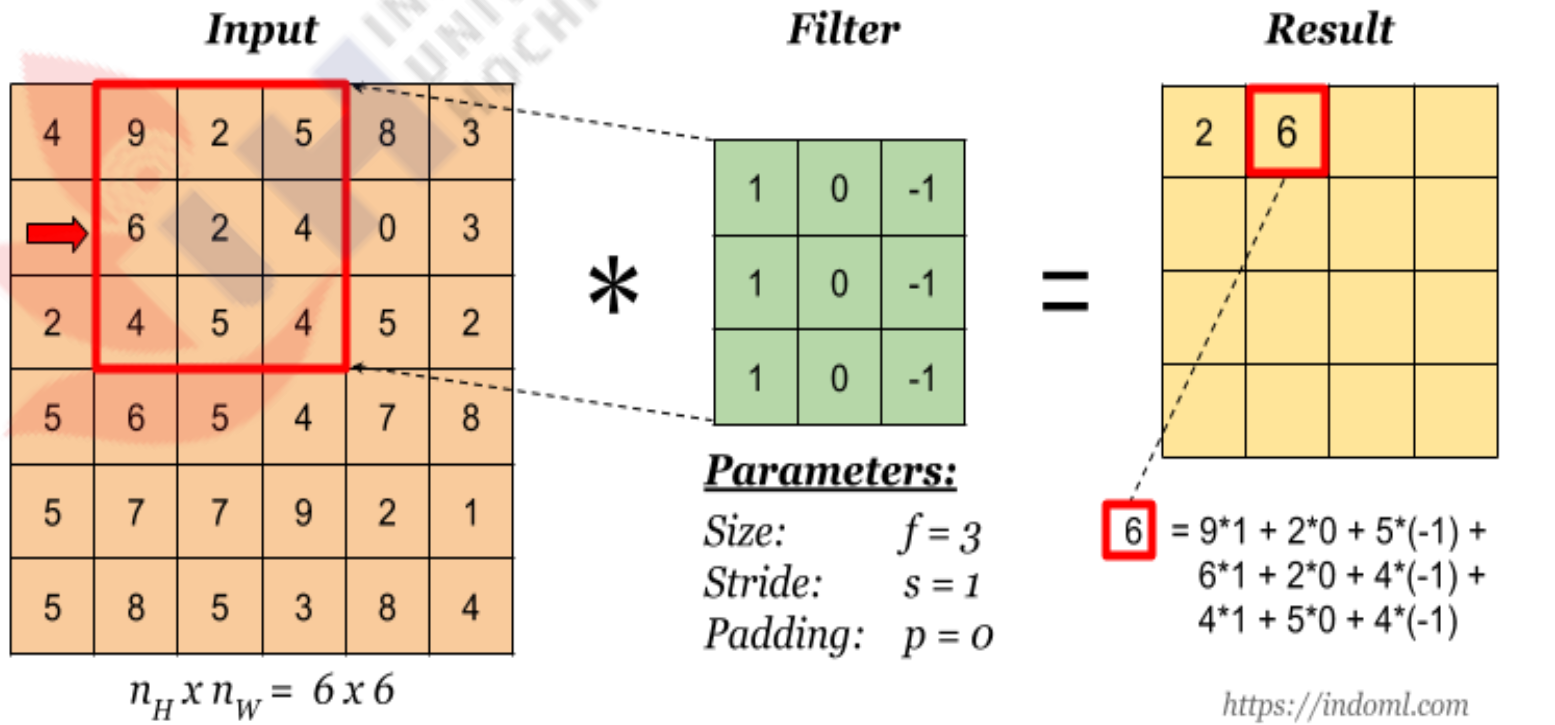
Input
Filter
Result
4 9 2 5 8 3
6 2 4 0 3
2 4 5 4 5 2
5 6 5 4 7 8
5 7 7 9 2 1
5 8 5 3 8 4
1 0 -1
1 0 -1
1 0 -1
2 6
Parameters:
Size: f = 3
Stride: s = 1
Padding: p = 0
6 = 9*1 + 2*0 + 5*(-1) +
6*1 + 2*0 + 4*(-1) +
4*1 + 5*0 + 4*(-1)
https://indoml.com

Stride

- Stride governs how many cells the filter is moved in the input to calculate the next cell in the result.

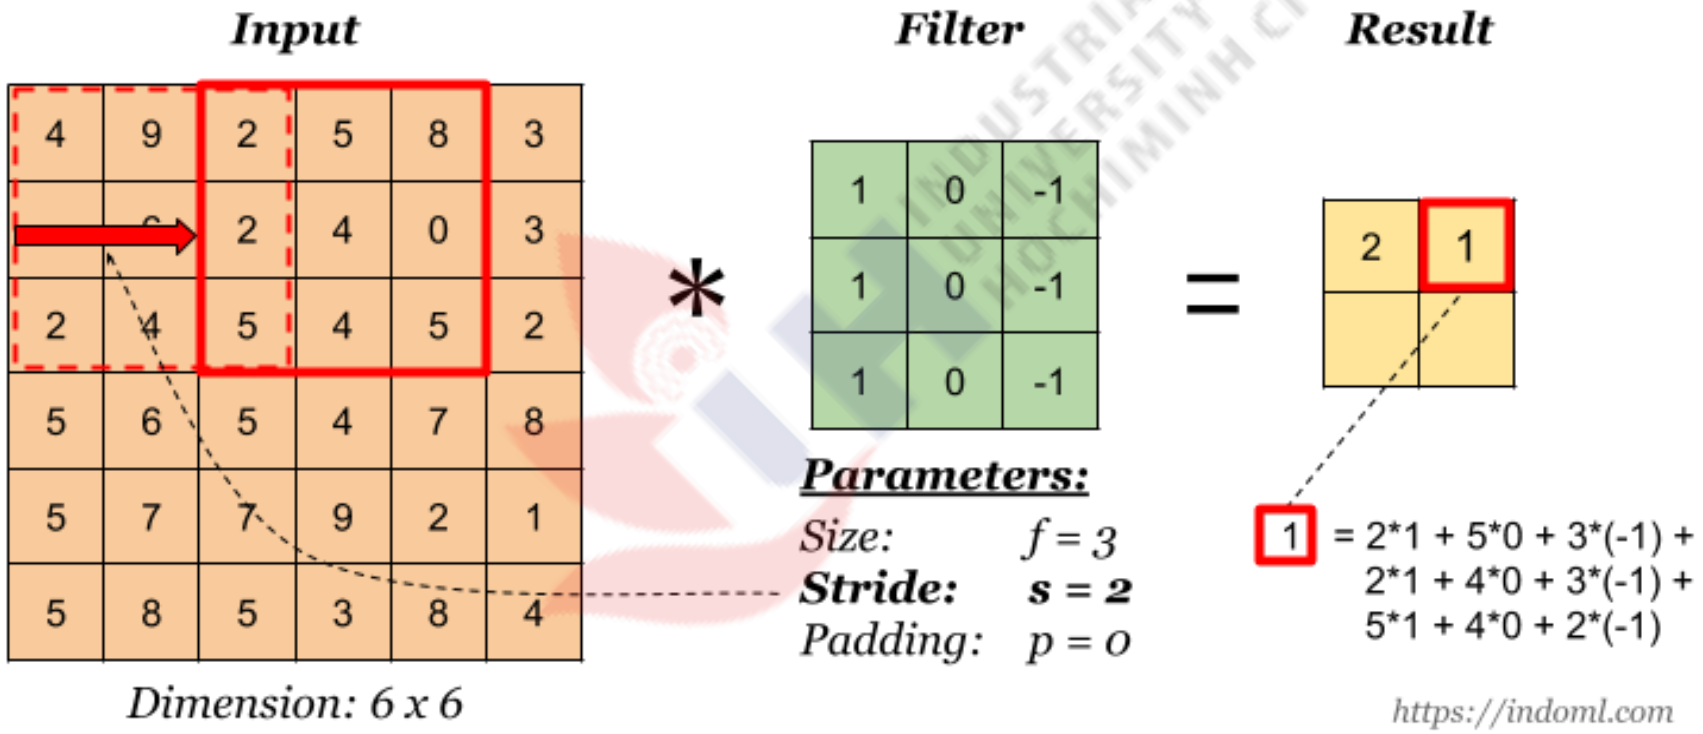

Padding

- important for building deeper networks
- keep more of the information at the border of an image, Without padding, very few values at the next layer would be affected by pixels as the edges of an image.

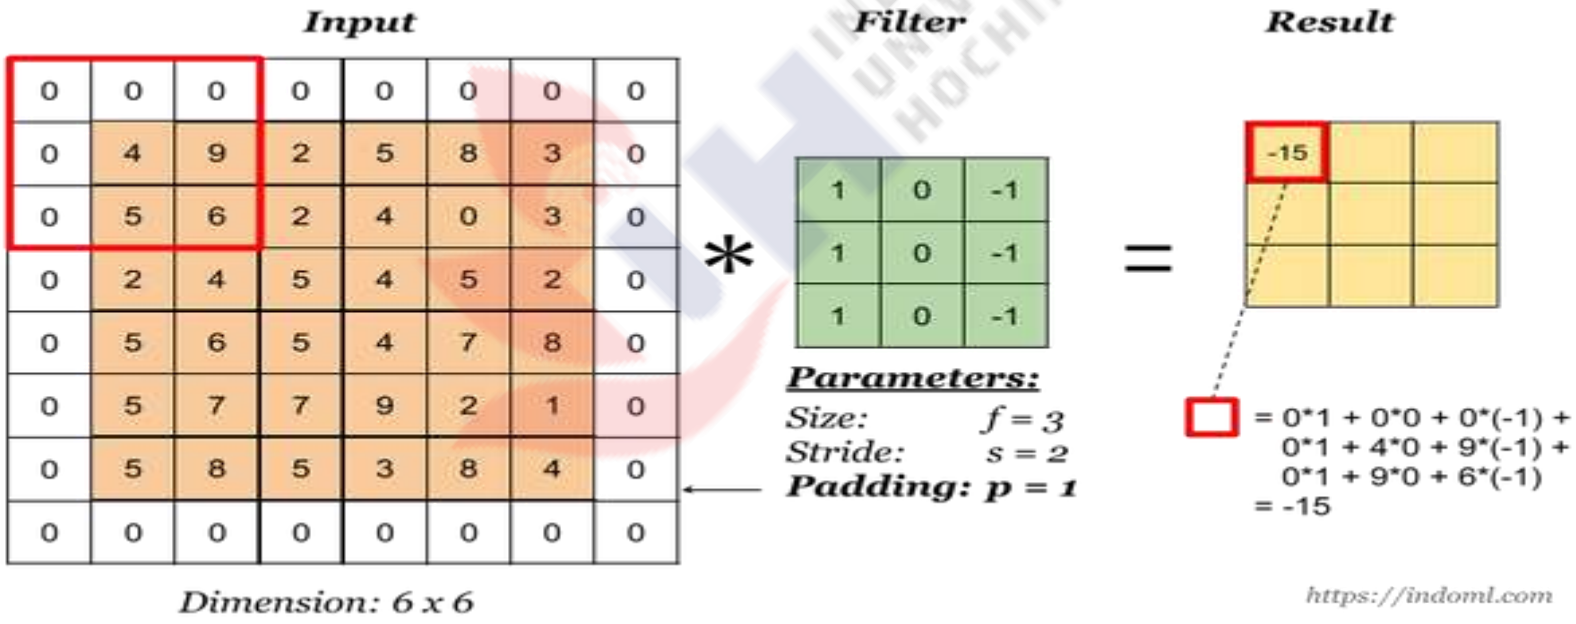

Convolution Operation on Volume

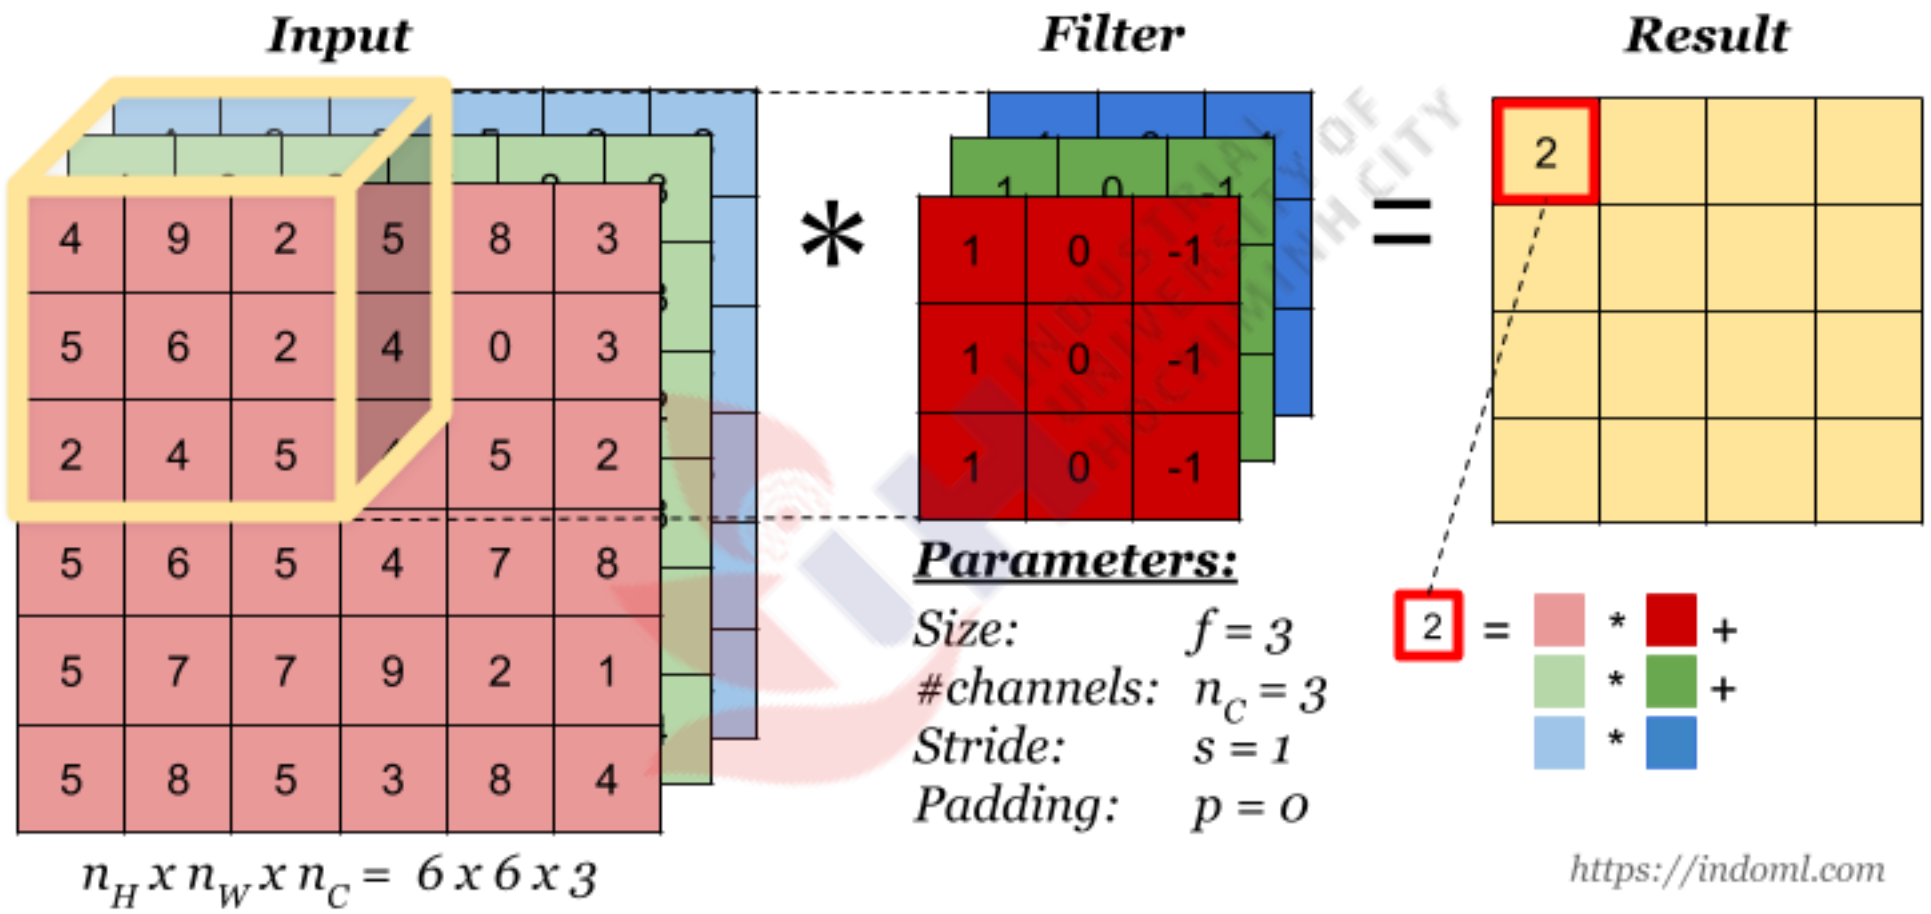

Convolution Operation with Multiple Filters

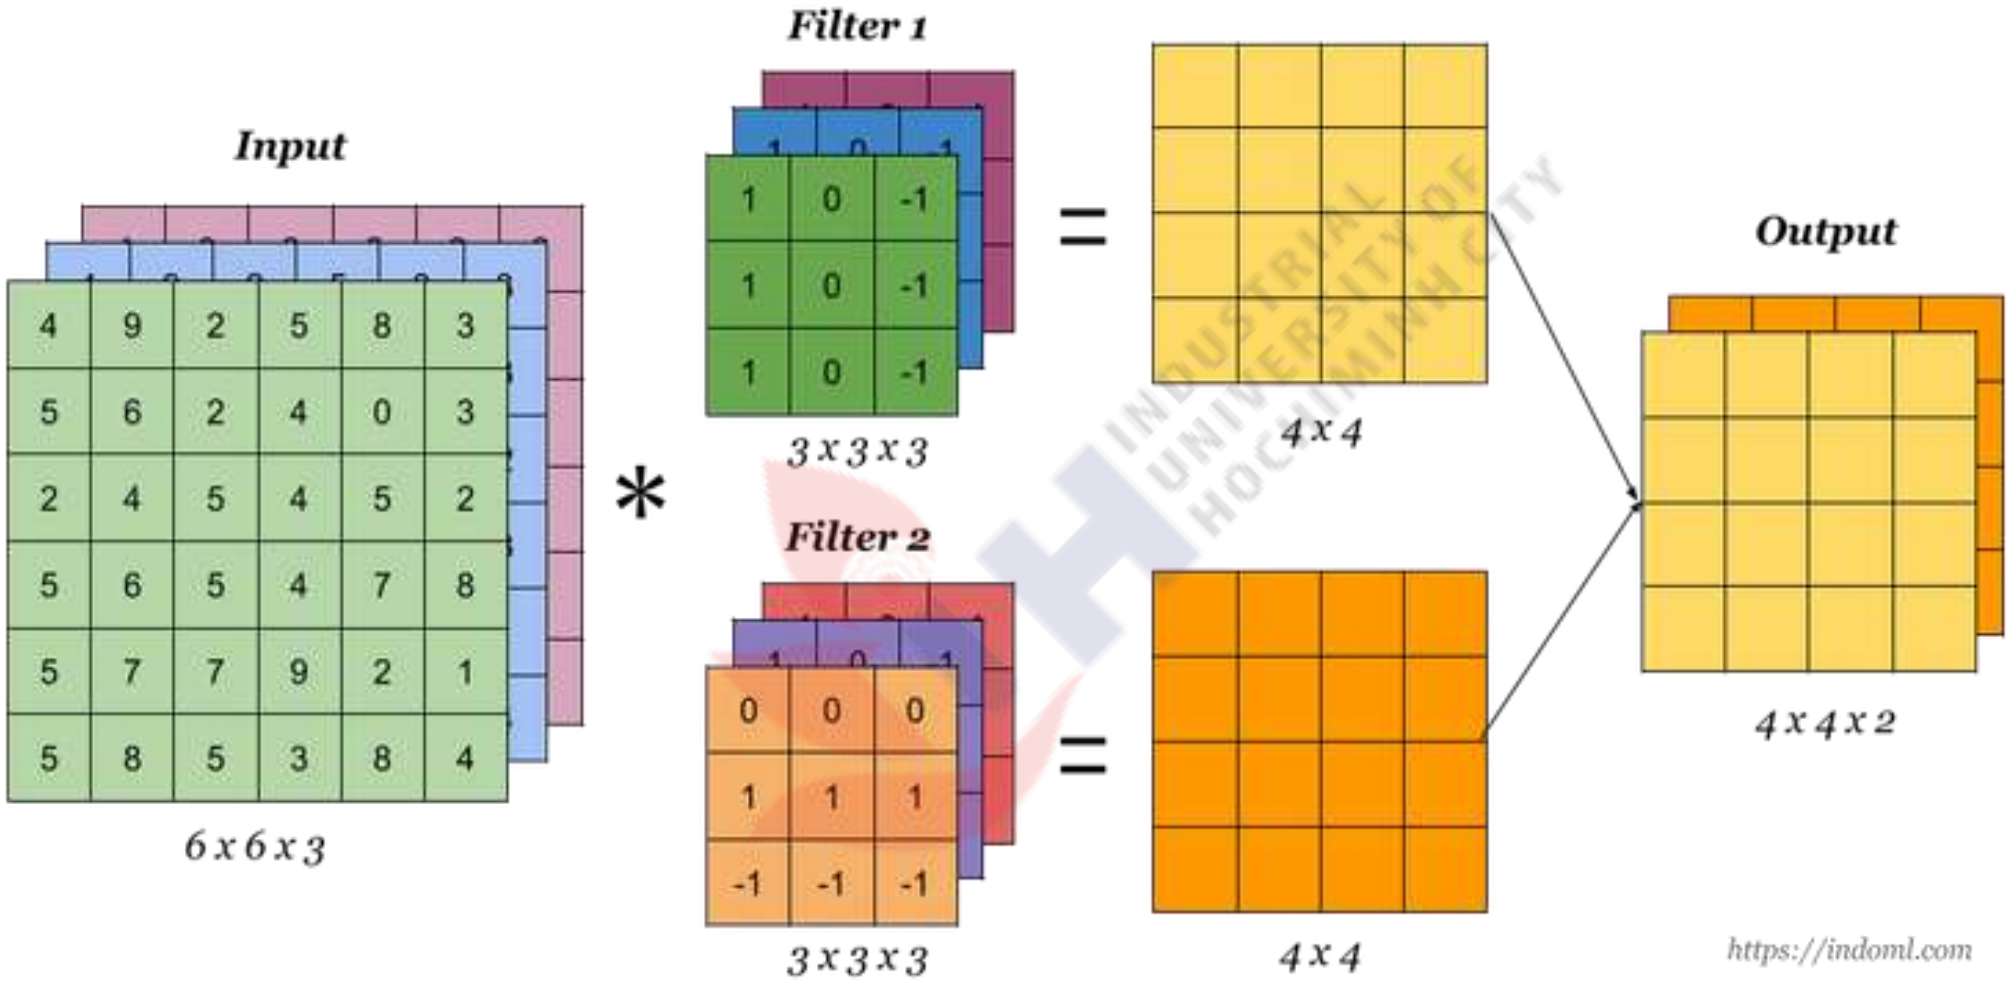

One Convolution Layer

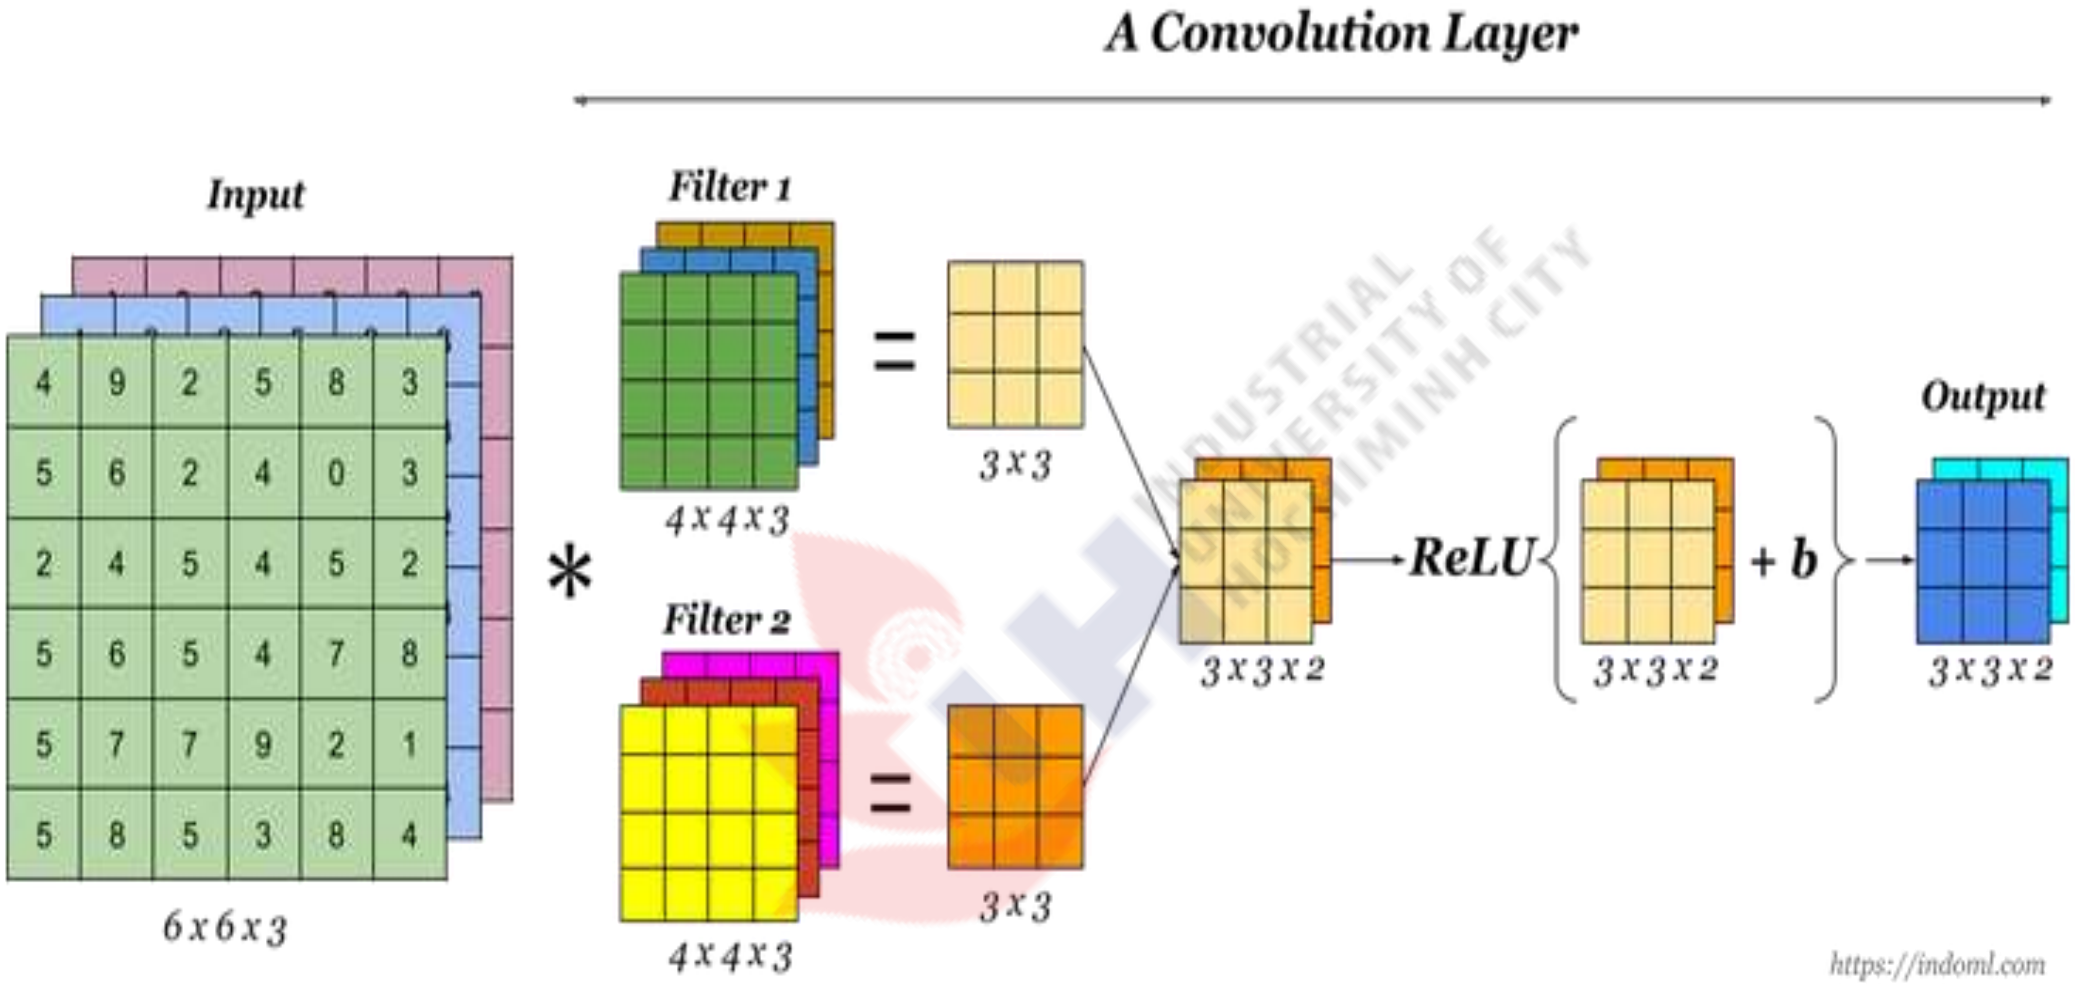

Sample Complete Network

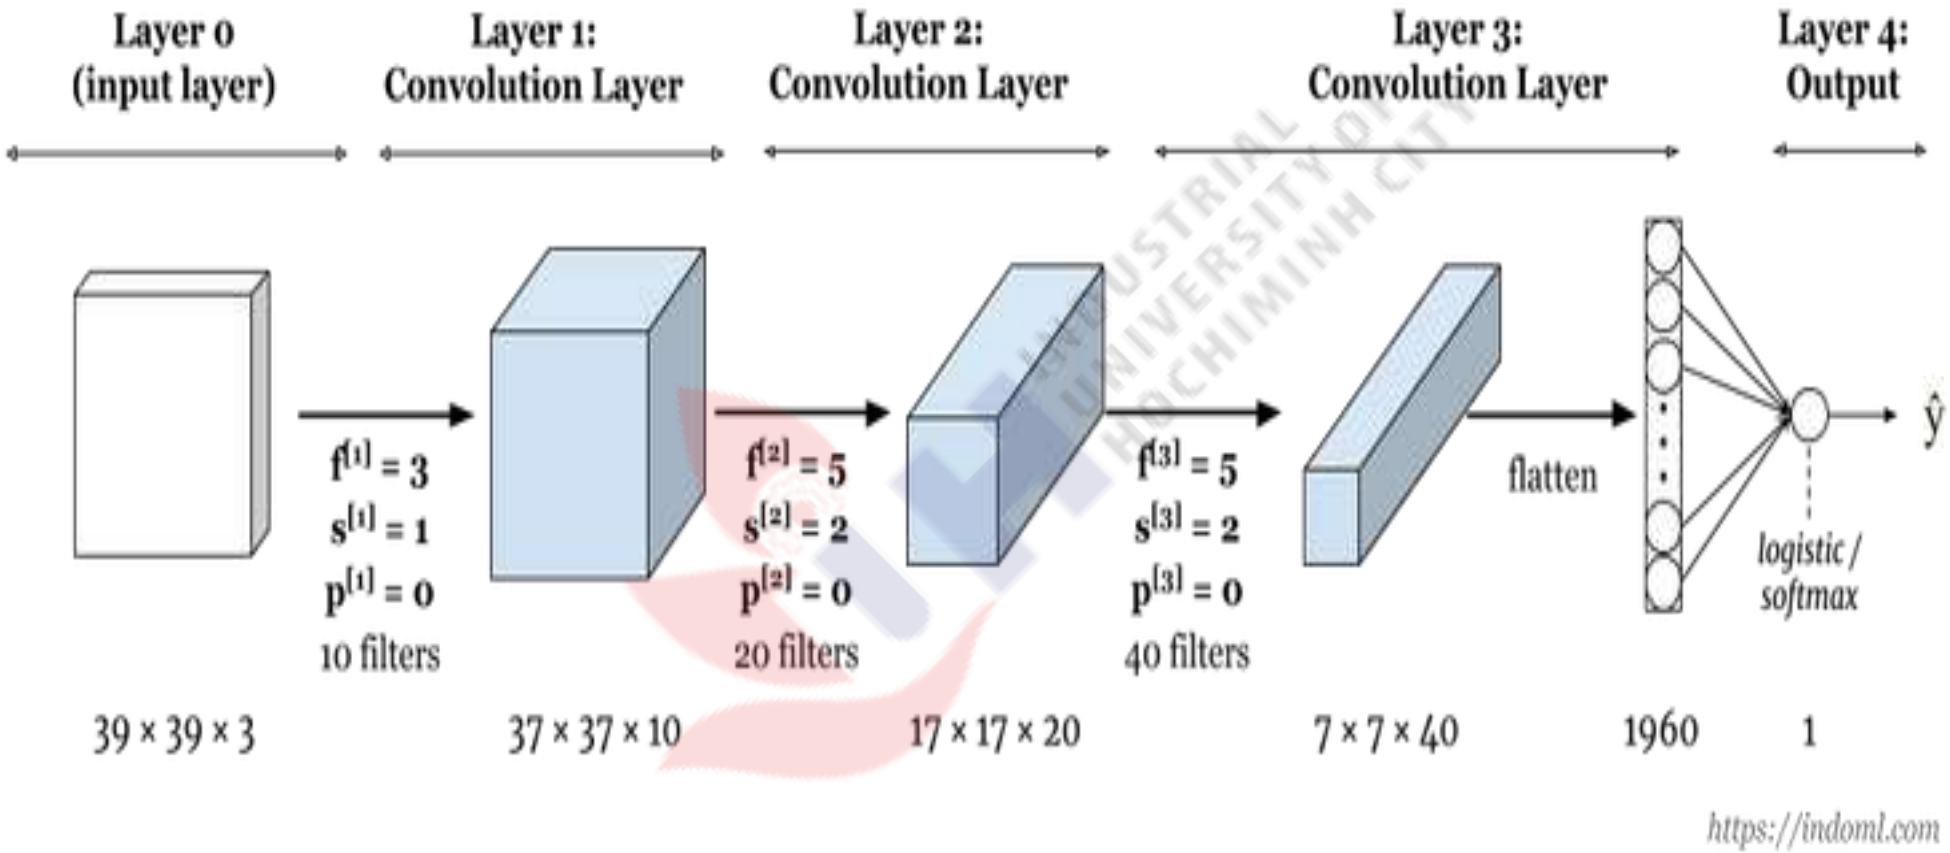

	0	1	0	
	1	-4	1	
	0	1	0	

0	0	0	0	0
0	1	1	1	0
0	1	1	1	0
0	1	1	1	0
0	0	0	0	0

Pooling Layer

- Pooling layer is used to reduce the size of the representations and to speed up calculations

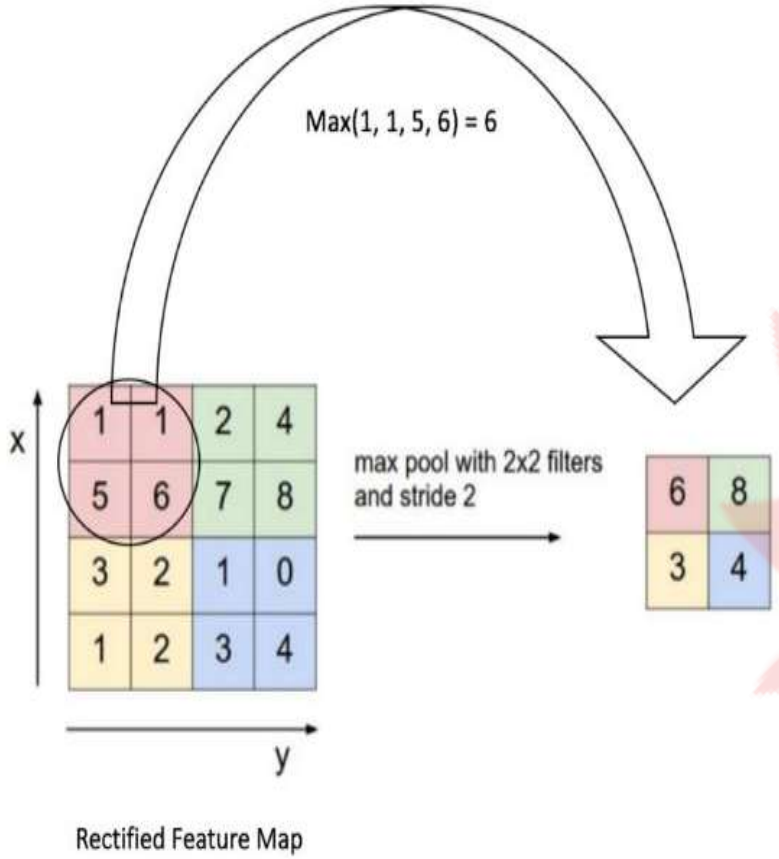

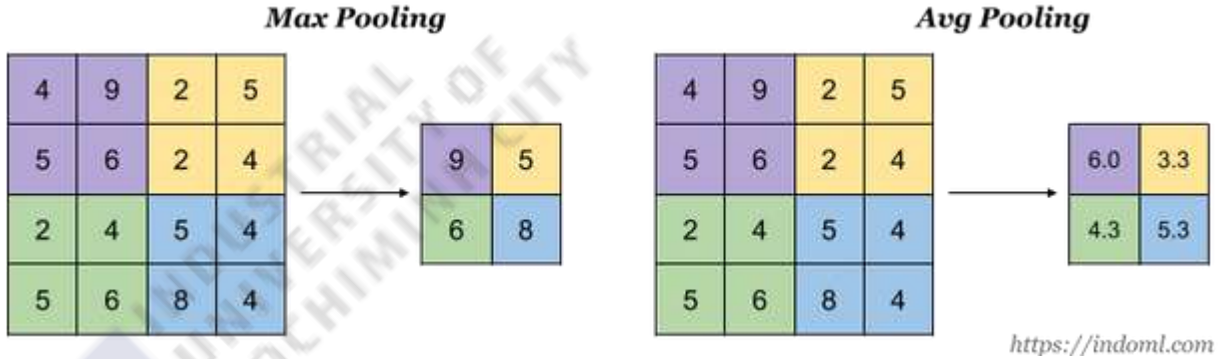

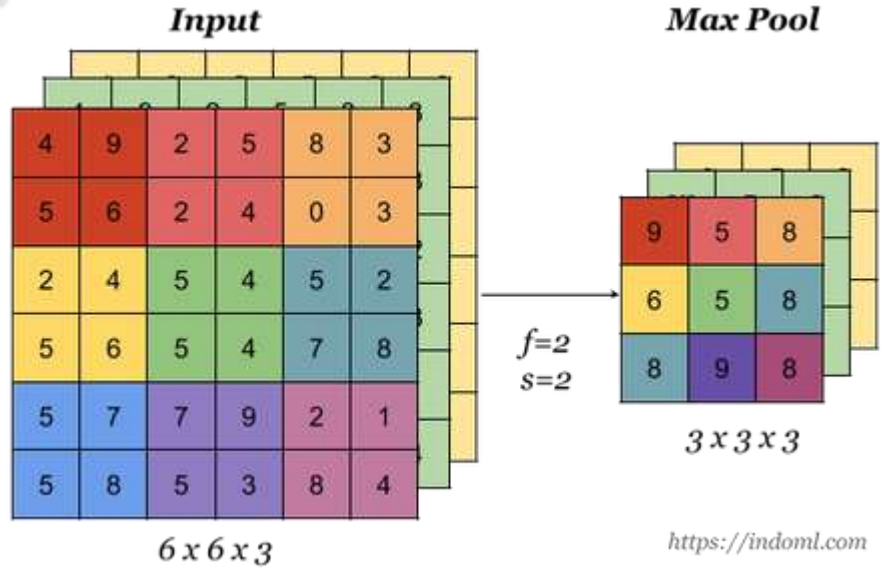

Flattening

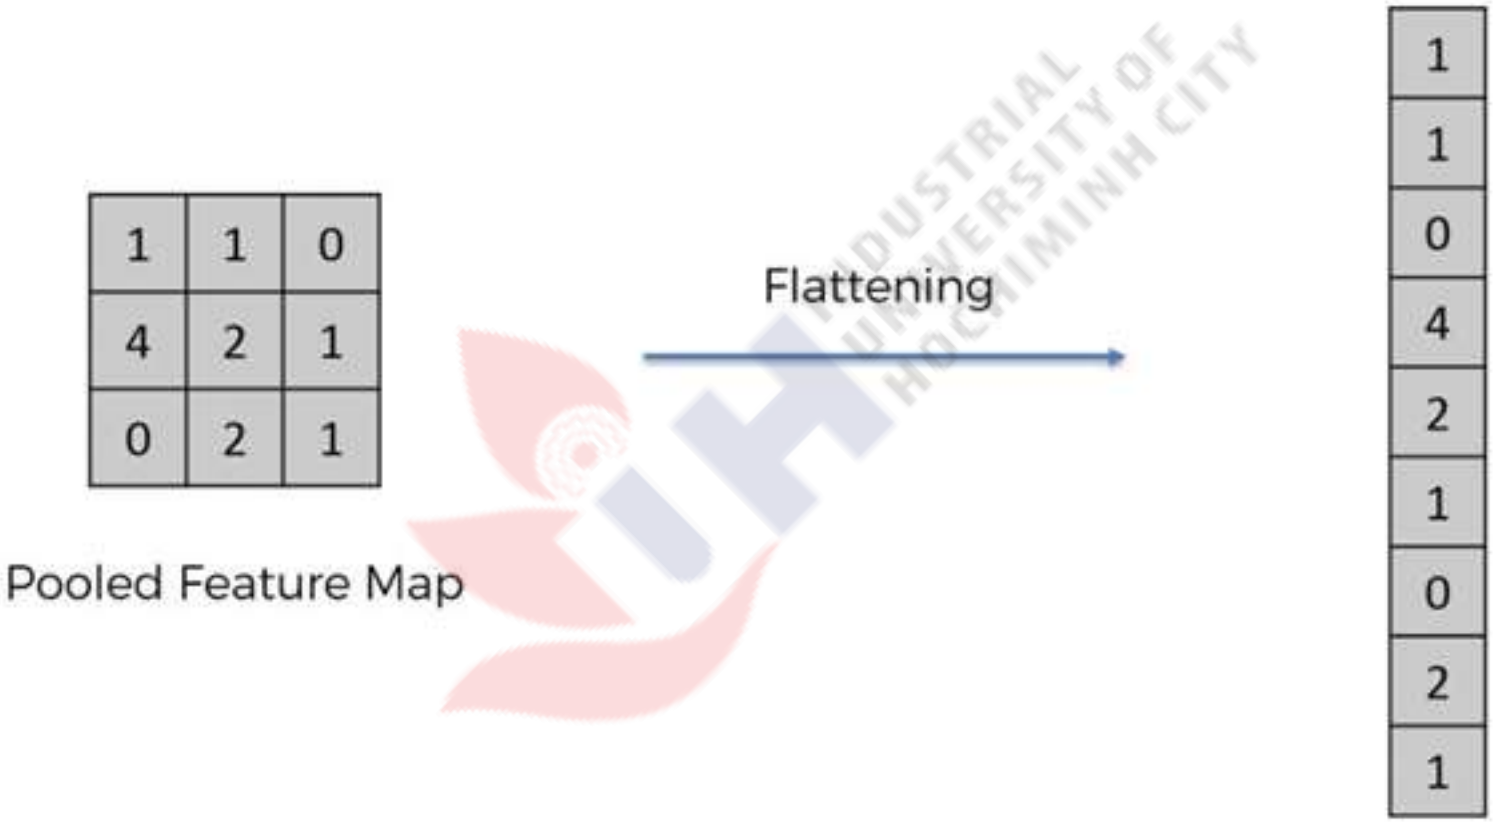

The process of building a CNN involves four major steps

- Convolution
- Pooling
- Flattening
- Full Connection

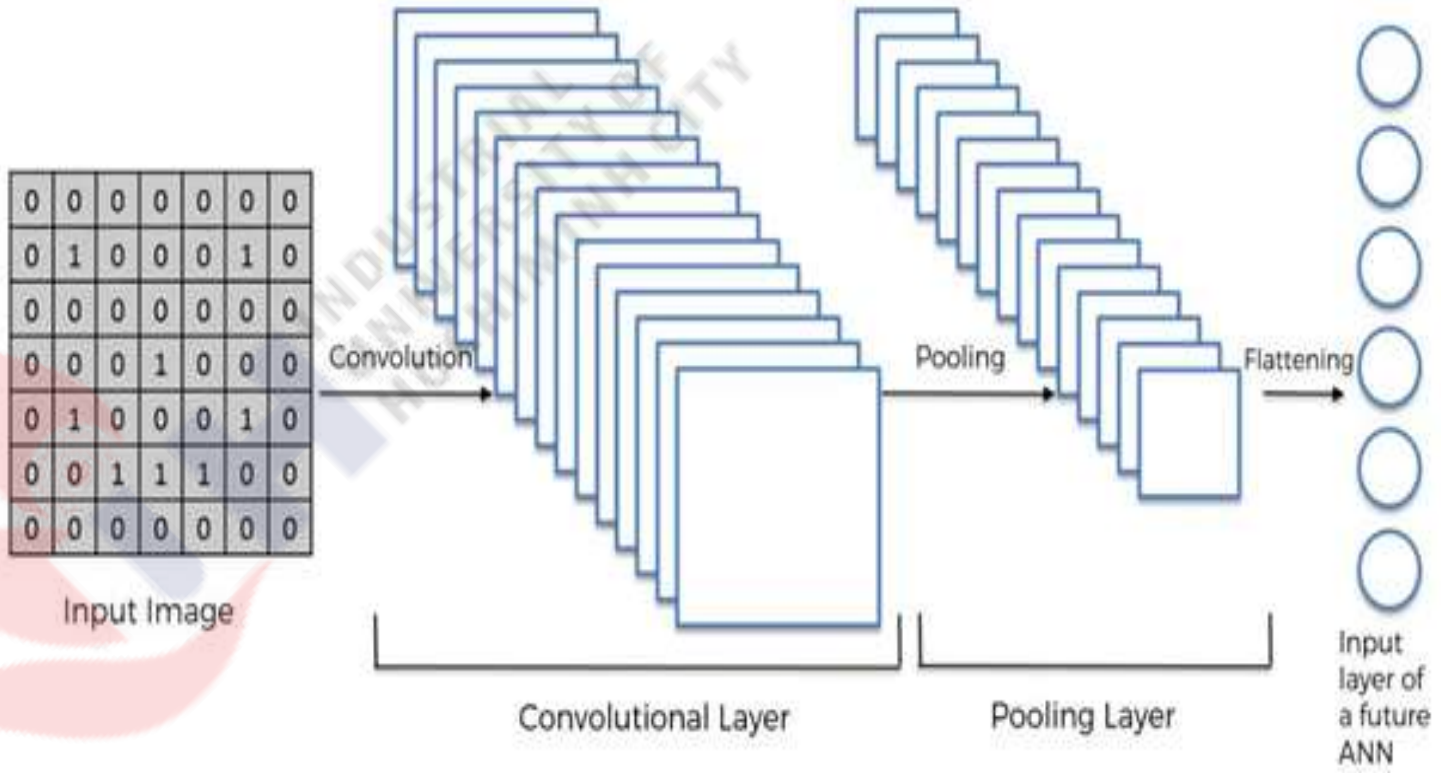

Full Connection

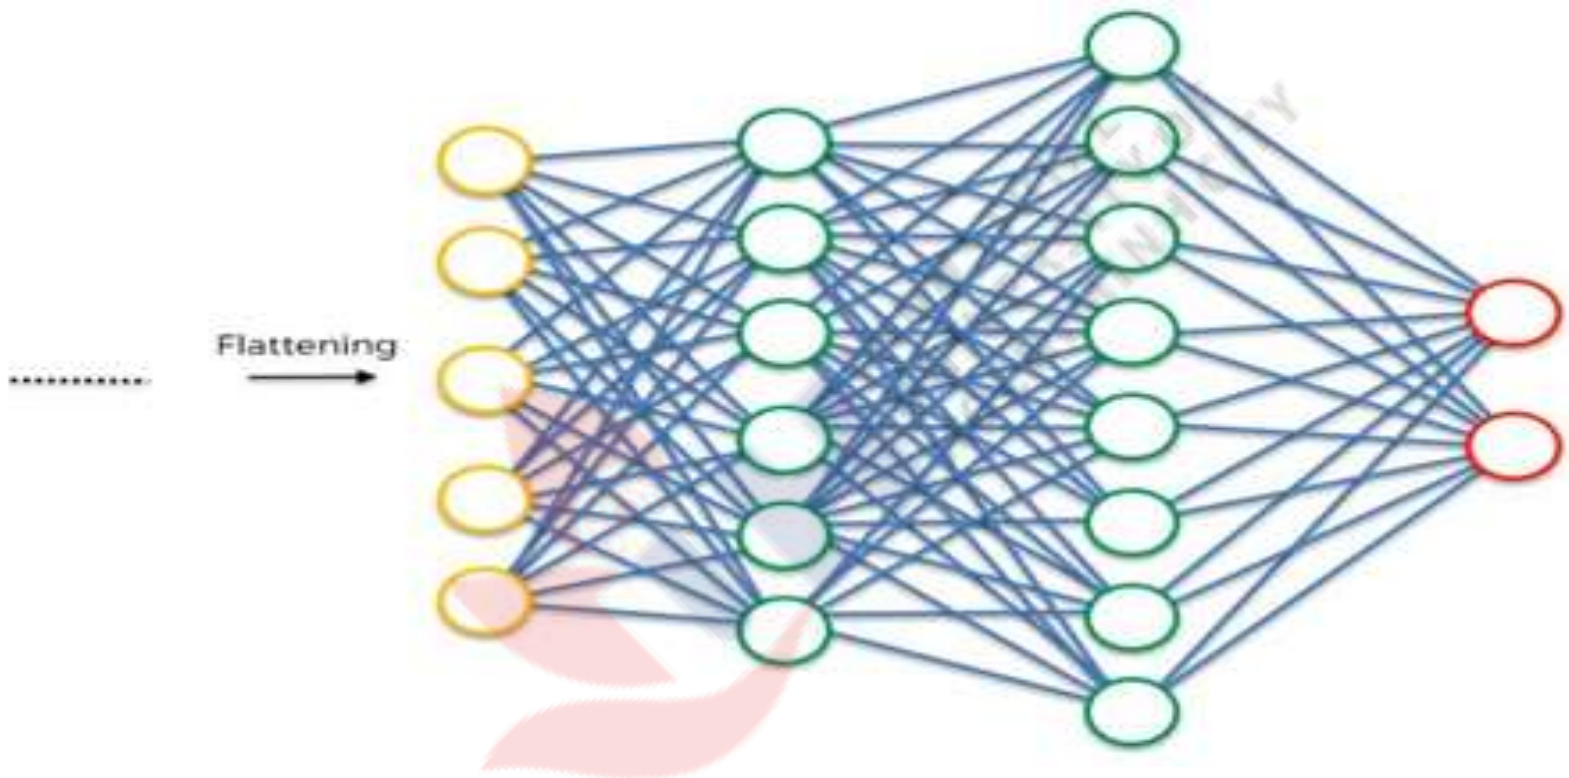

- **Putting it all together Training using Backpropagation**

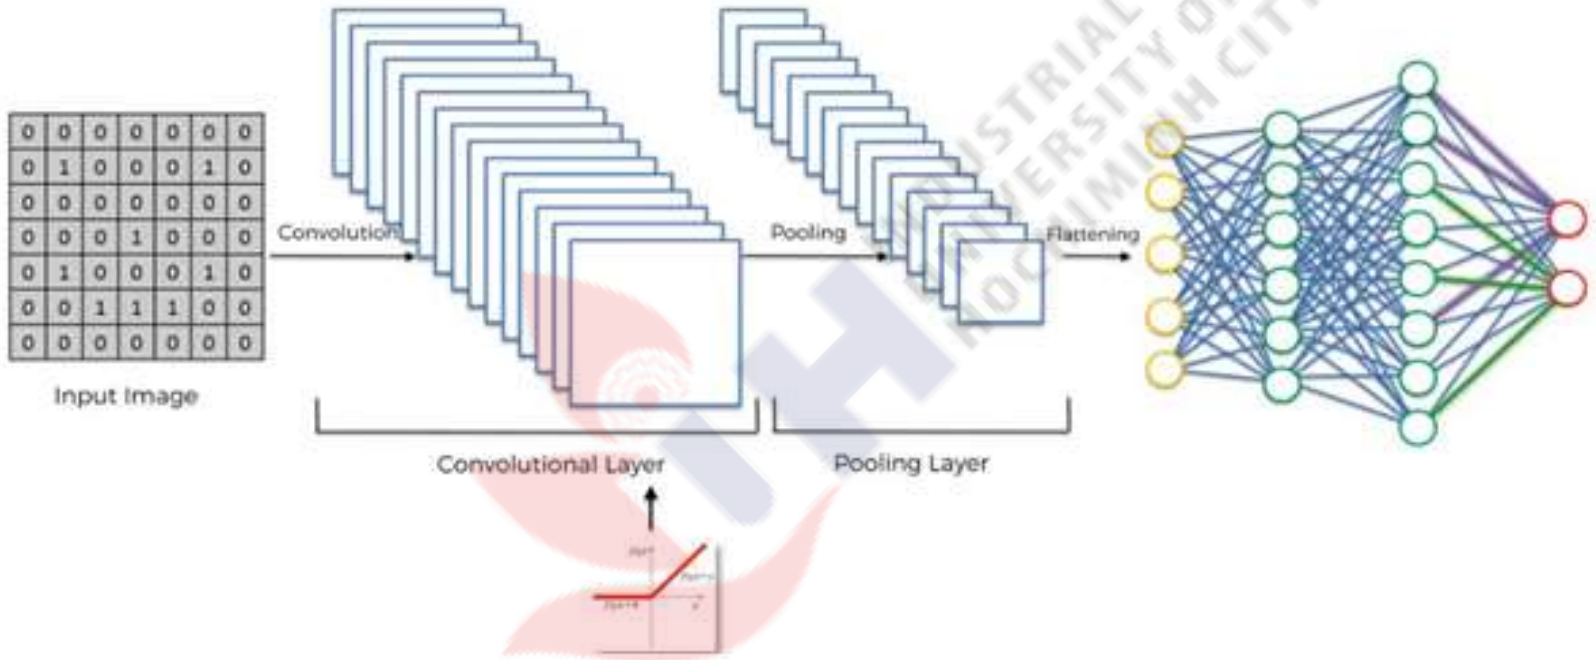

- Input Image = Boat
- Target Vector = [0, 0, 1, 0]

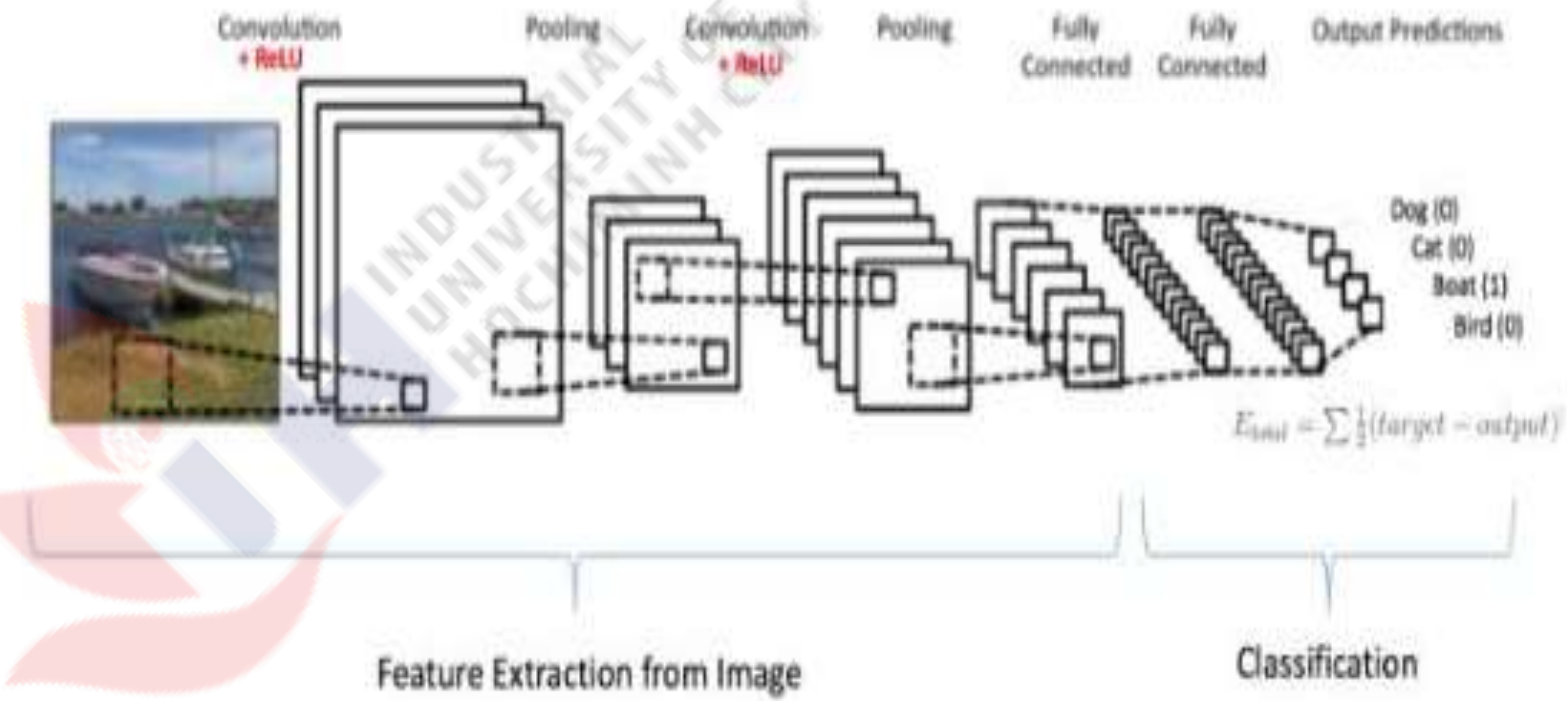

Training process

Step1: initialize all filters and parameters/weights with random values.

Step2: takes a training image as input => through the forward propagation step => finds the output probabilities for each class

Ex: output probabilities for the boat image above are [0.2, 0.4, 0.1, 0.3]

Step3: Calculate the total error at the output layer (summation over all 4 classes)

Total Error = $\sum$ ½ (target probability — output probability) ^{2}

Step4: Use Backpropagation to calculate the *gradients* of the error concerning all weights in the network and use *gradient descent* to update all filter values/weights and parameter values to minimize the output error.

Ex: [0.1, 0.1, 0.7, 0.1] is closer to the target vector [0, 0, 1, 0].

Parameters like the number of filters, filter sizes, the architecture of the network, etc. have all been fixed before Step 1 and do not change during the training process — only the values of the filter matrix and connection weights get updated.

Step5: Repeat steps 2–4 with all images in the training set.

Overfitting and dropout

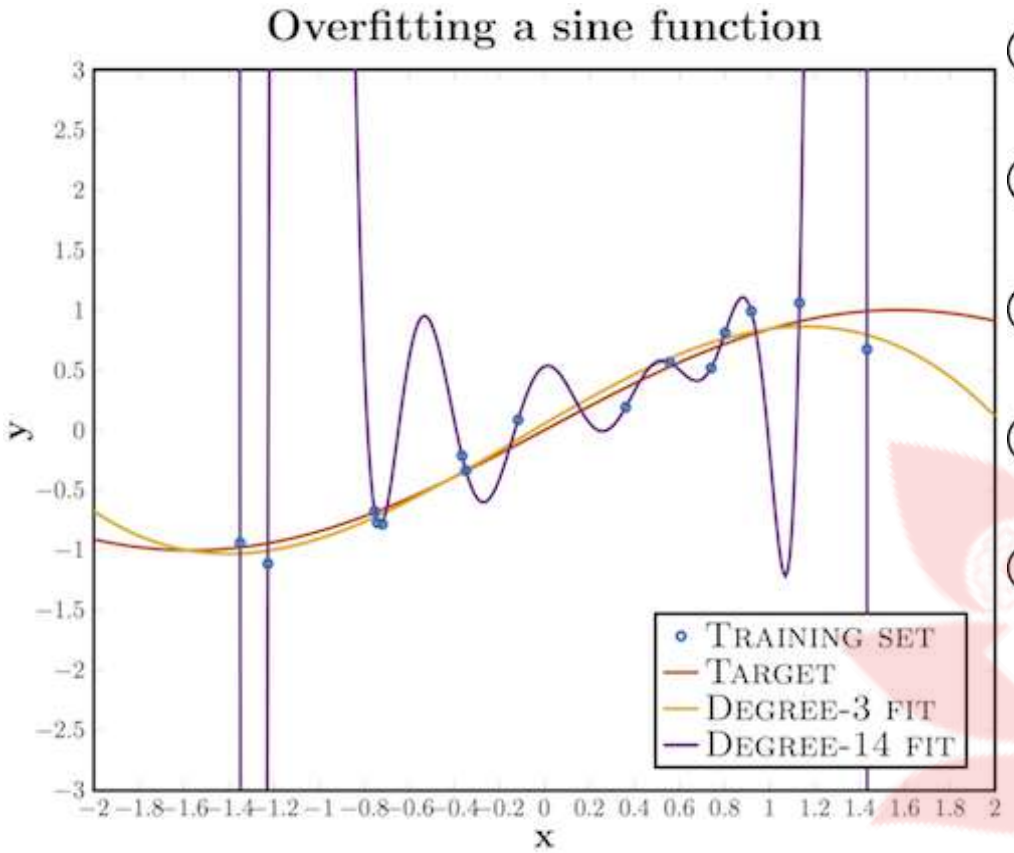

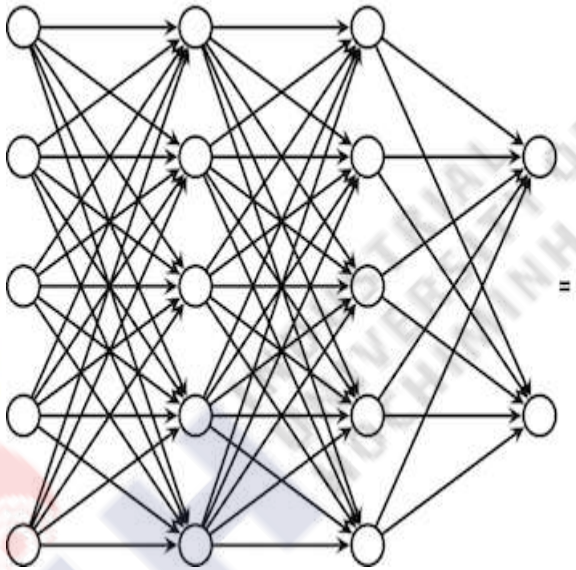

dropout

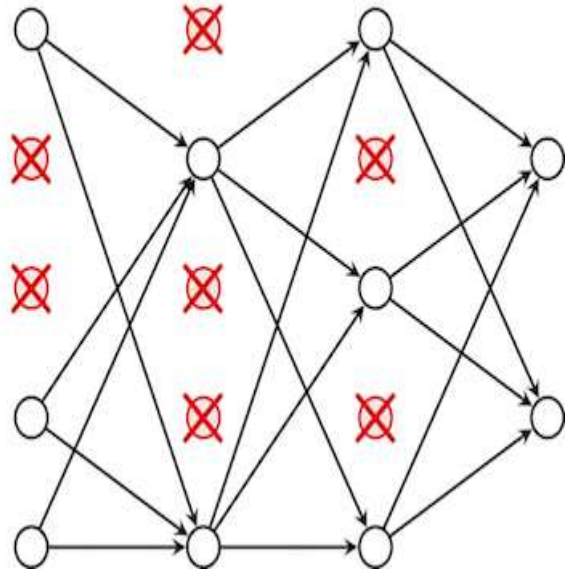

Well Known Architectures

- **LeNet – 5**
-

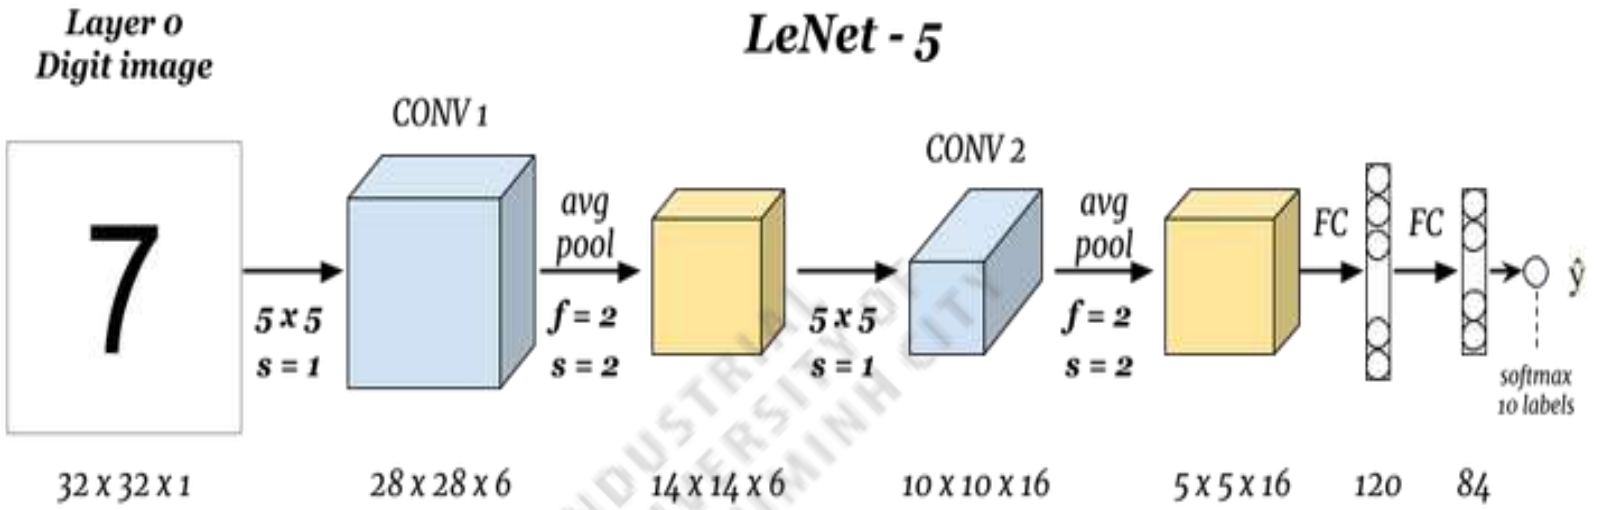

Layer	Layer Type	Feature Maps	Size	Kernel Size	Stride	Activation
Input	Image	1	32x32	-	-	-
1	Convolution	6	28x82	5x5	1	tanh
2	Average Pooling	6	14x14	2x2	2	tanh
3	Convolution	16	10x10	5x5	1	tanh
4	Average Pooling	16	5x5	2x2	2	tanh
5	Convolution	120	1x1	5x5	1	tanh
6	Fully Connected	-	84	-	-	tanh
Output	Fully Connected	-	10	-	-	softmax

Classic Network: AlexNet

Layer		Feature Map	Size	Kernel Size	Stride	Activation
Input	Image	1	227x227x3	-	-	-
1	Convolution	96	55 x 55 x 96	11x11	4	relu
	Max Pooling	96	27 x 27 x 96	3x3	2	relu
2	Convolution	256	27 x 27 x 256	5x5	1	relu
	Max Pooling	256	13 x 13 x 256	3x3	2	relu
3	Convolution	384	13 x 13 x 384	3x3	1	relu
4	Convolution	384	13 x 13 x 384	3x3	1	relu
5	Convolution	256	13 x 13 x 256	3x3	1	relu
	Max Pooling	256	6 x 6 x 256	3x3	2	relu
6	FC	-	9216	-	-	relu
7	FC	-	4096	-	-	relu
8	FC	-	4096	-	-	relu
Output	FC	-	1000	-	-	Softmax

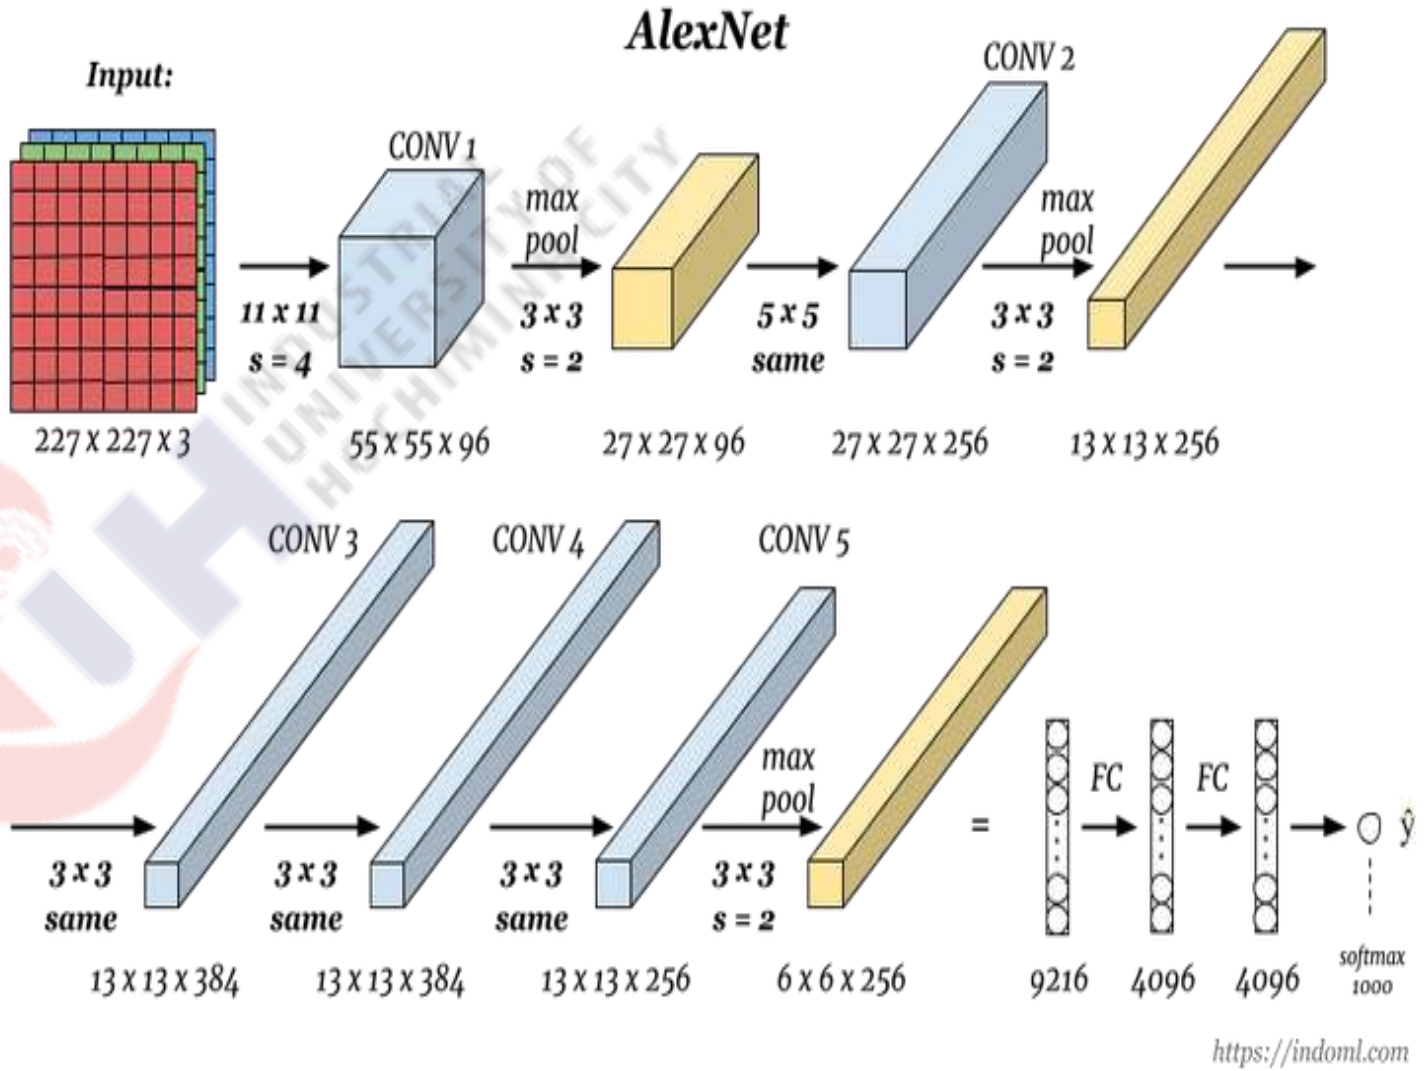

Classic Network: VGG-16

	Layer	Feature Map	Size	Kernel Size	Stride	Activation
Input	Image	1	224 x 224 x 3	-	-	-
1	2 X Convolution	64	224 x 224 x 64	3x3	1	relu
	Max Pooling	64	112 x 112 x 64	3x3	2	relu
3	2 X Convolution	128	112 x 112 x 128	3x3	1	relu
	Max Pooling	128	56 x 56 x 128	3x3	2	relu
5	2 X Convolution	256	56 x 56 x 256	3x3	1	relu
	Max Pooling	256	28 x 28 x 256	3x3	2	relu
7	3 X Convolution	512	28 x 28 x 512	3x3	1	relu
	Max Pooling	512	14 x 14 x 512	3x3	2	relu
10	3 X Convolution	512	14 x 14 x 512	3x3	1	relu
	Max Pooling	512	7 x 7 x 512	3x3	2	relu
13	FC	-	25088	-	-	relu
14	FC	-	4096	-	-	relu
15	FC	-	4096	-	-	relu
Output	FC	-	1000	-	-	Softmax

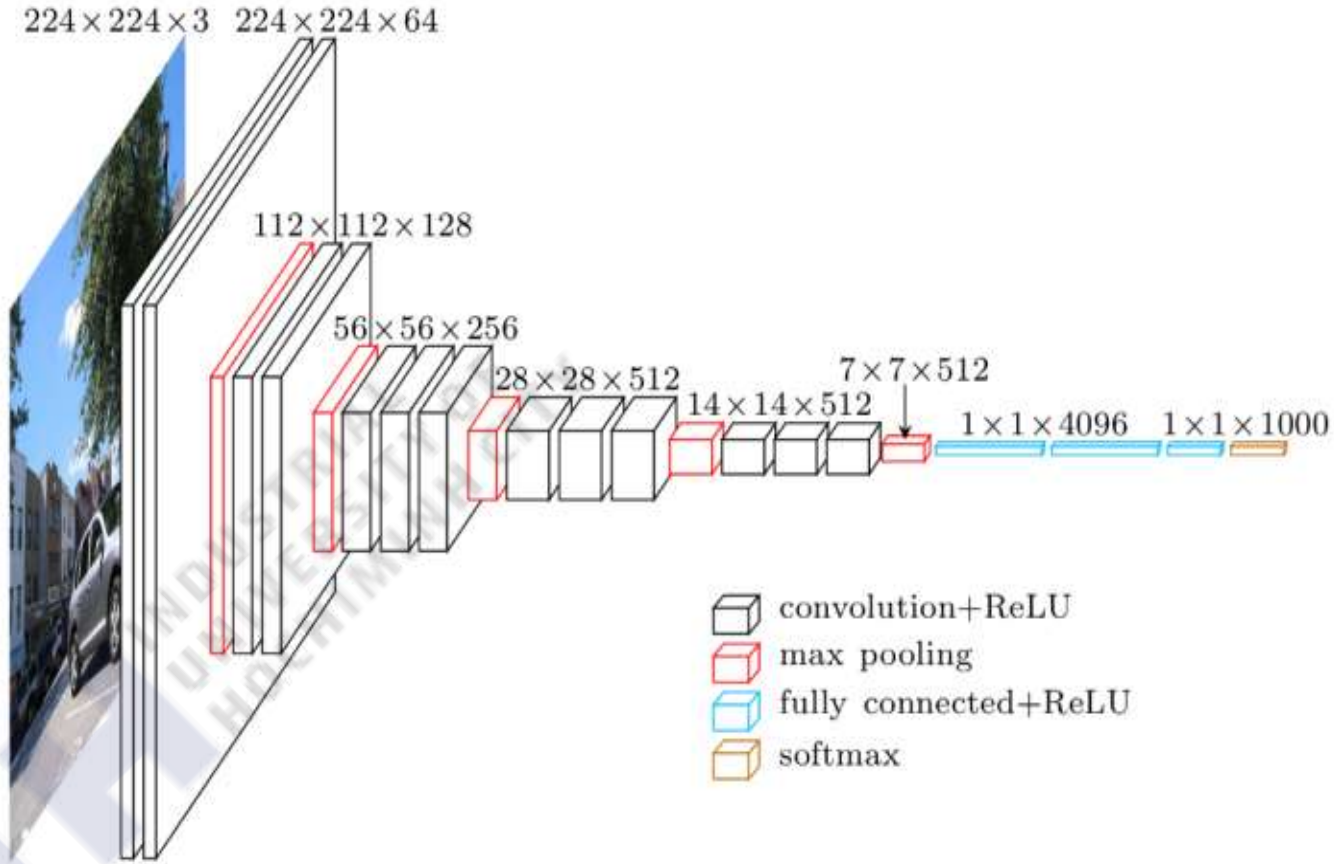

ResNet

- Vấn đề trong mạng deep neural networks sẽ khó huấn luyện khi đạt đến 1 cố layer nào đó, training error sẽ tăng trở lại, ngoài ra còn gặp vấn đề triệt tiêu hoặc đạo hàm quá lớn(exploding and vanishing gradients problem)
- Giải quyết vấn đề:
- $z^{[l+2]} = W^{[l+2]} a^{[l+1]} + b^{[l+2]}$
- $a^{[l+2]} = g^{[l+2]}(z^{[l+2]} + a^{[l]})$

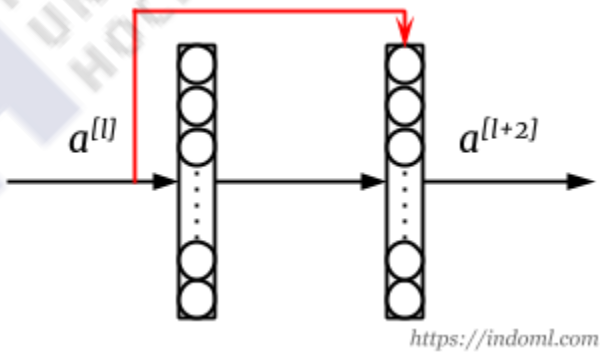

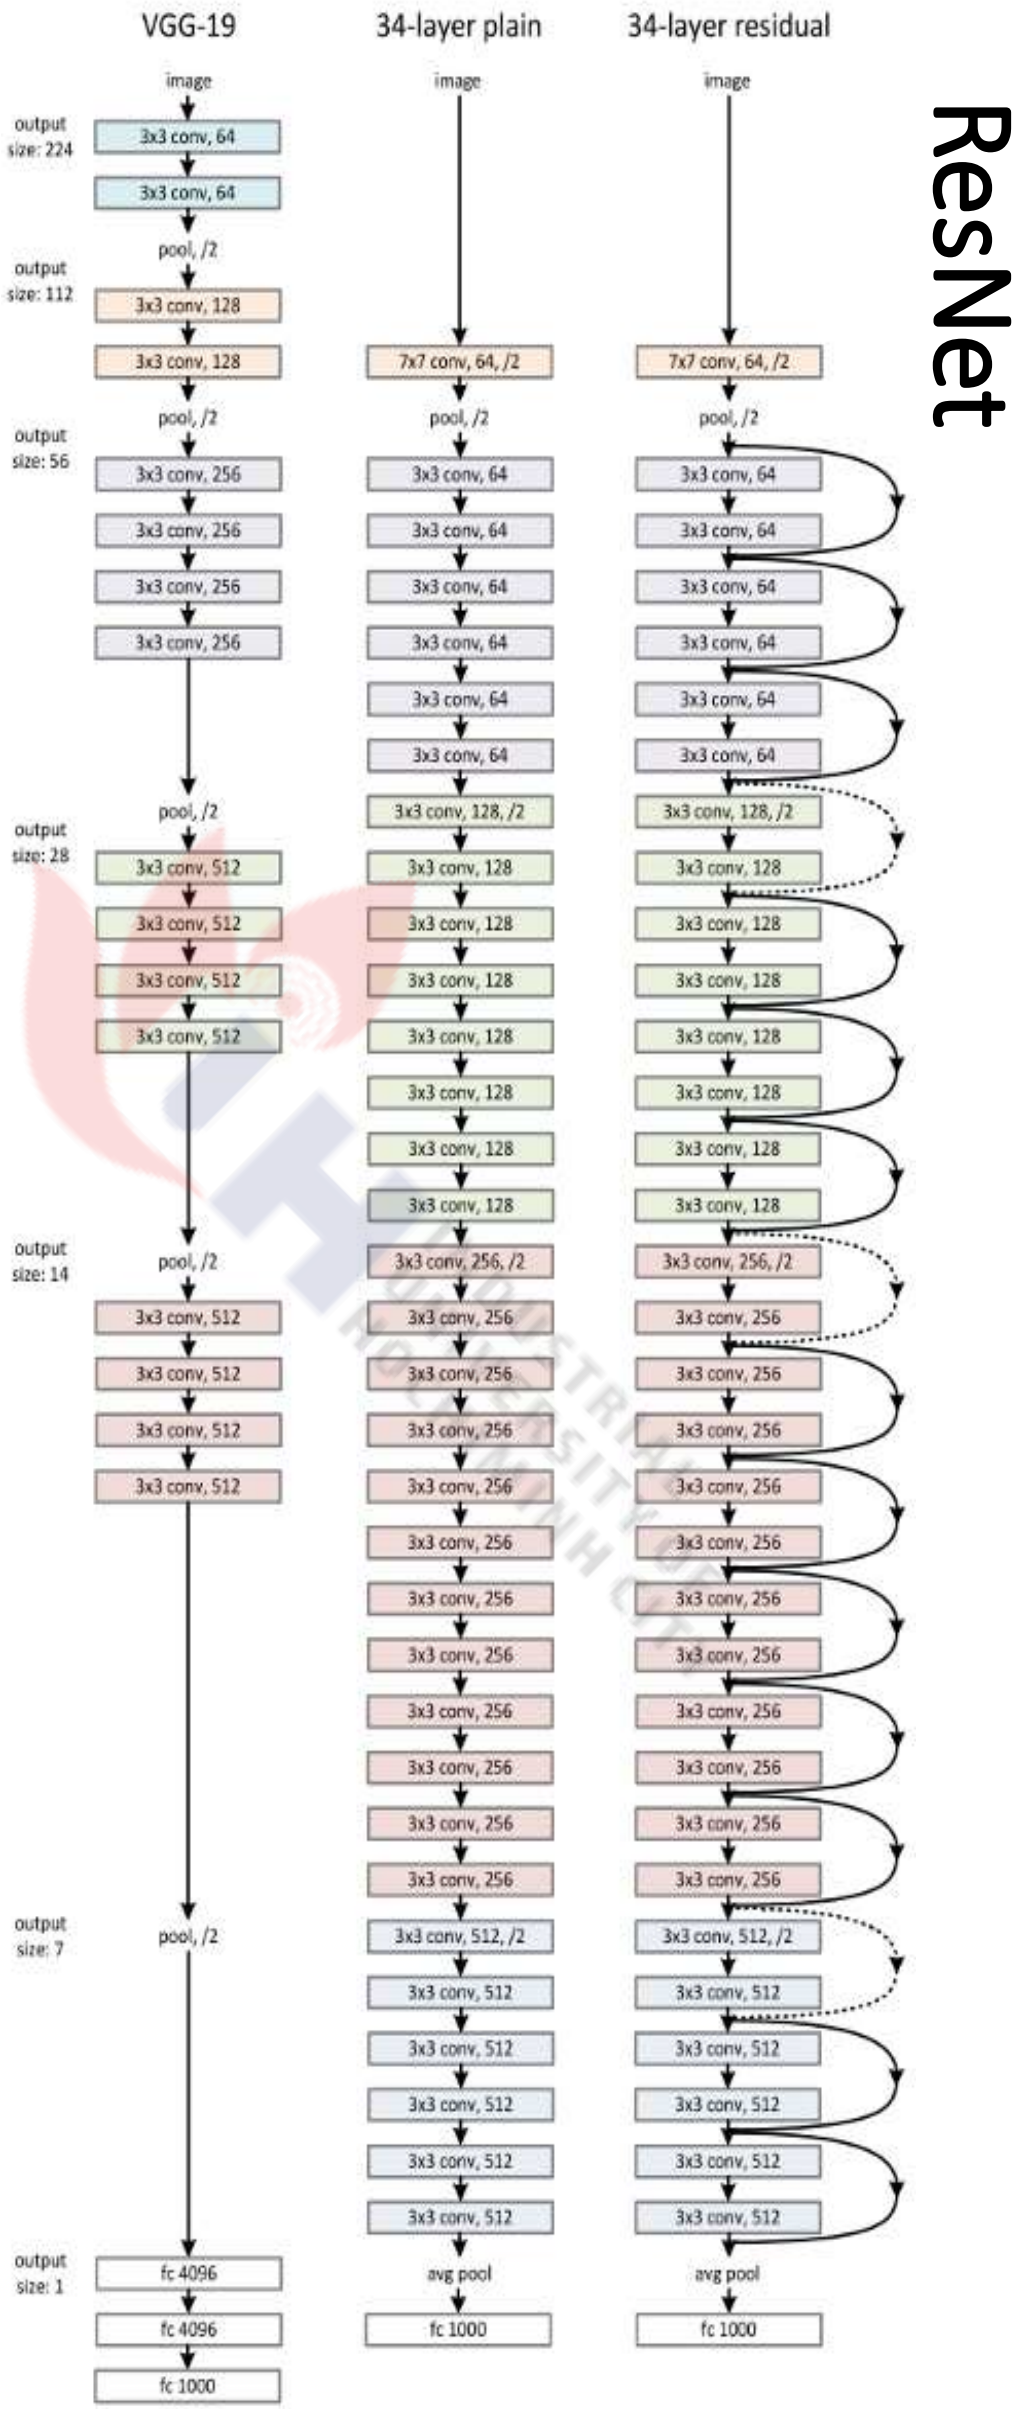
VGG-19
34-layer plain
34-layer residual
ResNet
image
output size: 224
3x3 conv, 64
3x3 conv, 64
pool, /2
output size: 112
3x3 conv, 128
3x3 conv, 128
pool, /2
output size: 56
3x3 conv, 256
3x3 conv, 256
3x3 conv, 256
3x3 conv, 256
pool, /2
output size: 28
3x3 conv, 512
3x3 conv, 512
3x3 conv, 512
3x3 conv, 512
pool, /2
output size: 14
3x3 conv, 512
3x3 conv, 512
3x3 conv, 512
3x3 conv, 512
pool, /2
output size: 7
fc 4096
fc 4096
fc 1000
output size: 1
7x7 conv, 64, /2
pool, /2
3x3 conv, 64
3x3 conv, 128, /2
3x3 conv, 128
3x3 conv, 256, /2
3x3 conv, 256
3x3 conv, 512, /2
3x3 conv, 512
avg pool
fc 1000

Inception

- Thay vì chọn kích thước filter 1 cách thủ công, network quyết định chọn cái tốt nhất để đưa vào layer
- **GoogLeNet**,

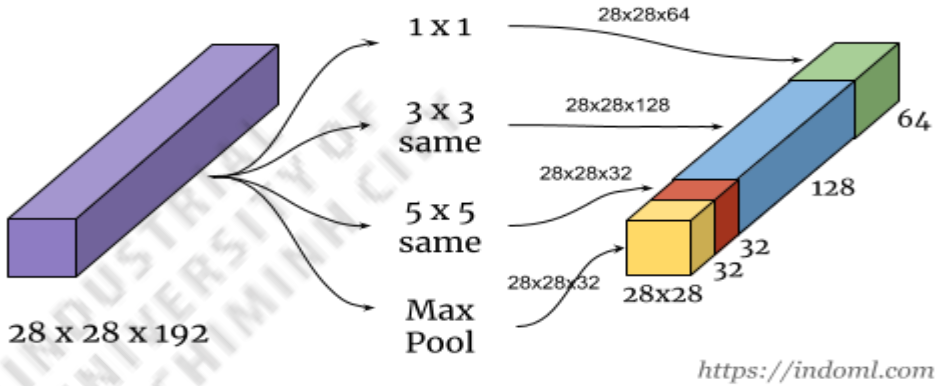

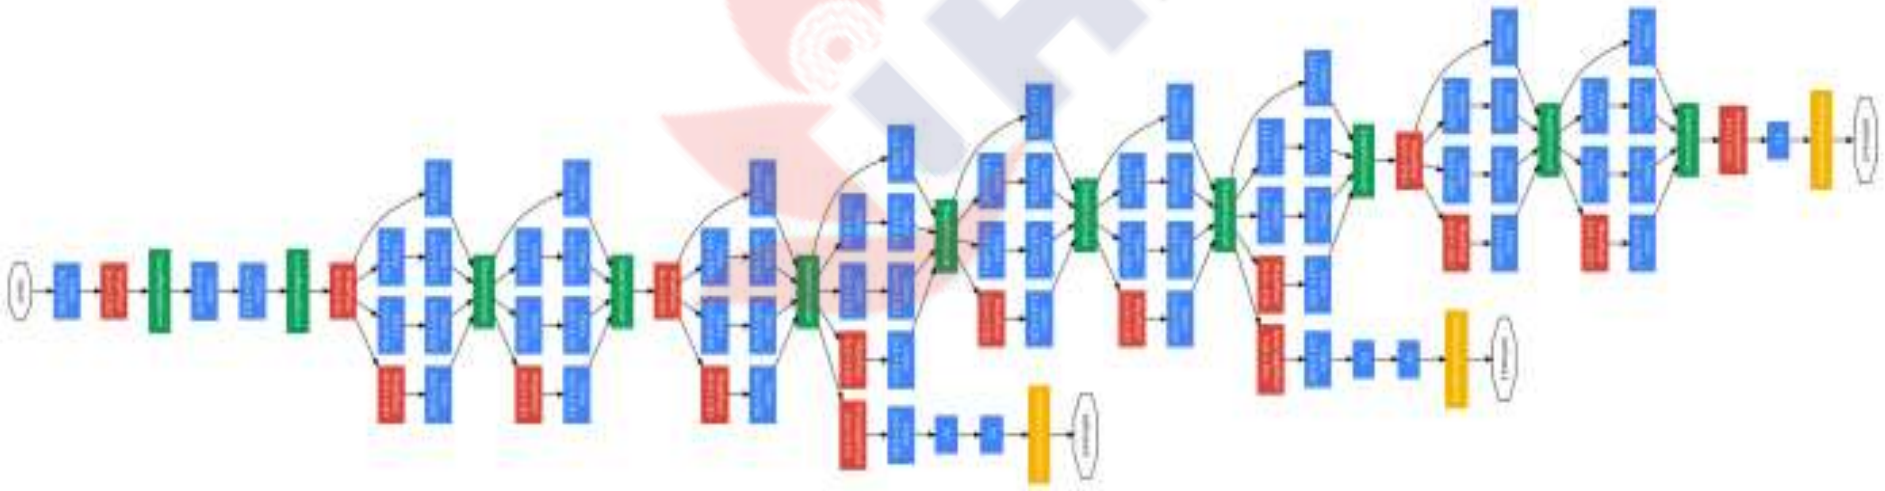

Retrain ResNet50

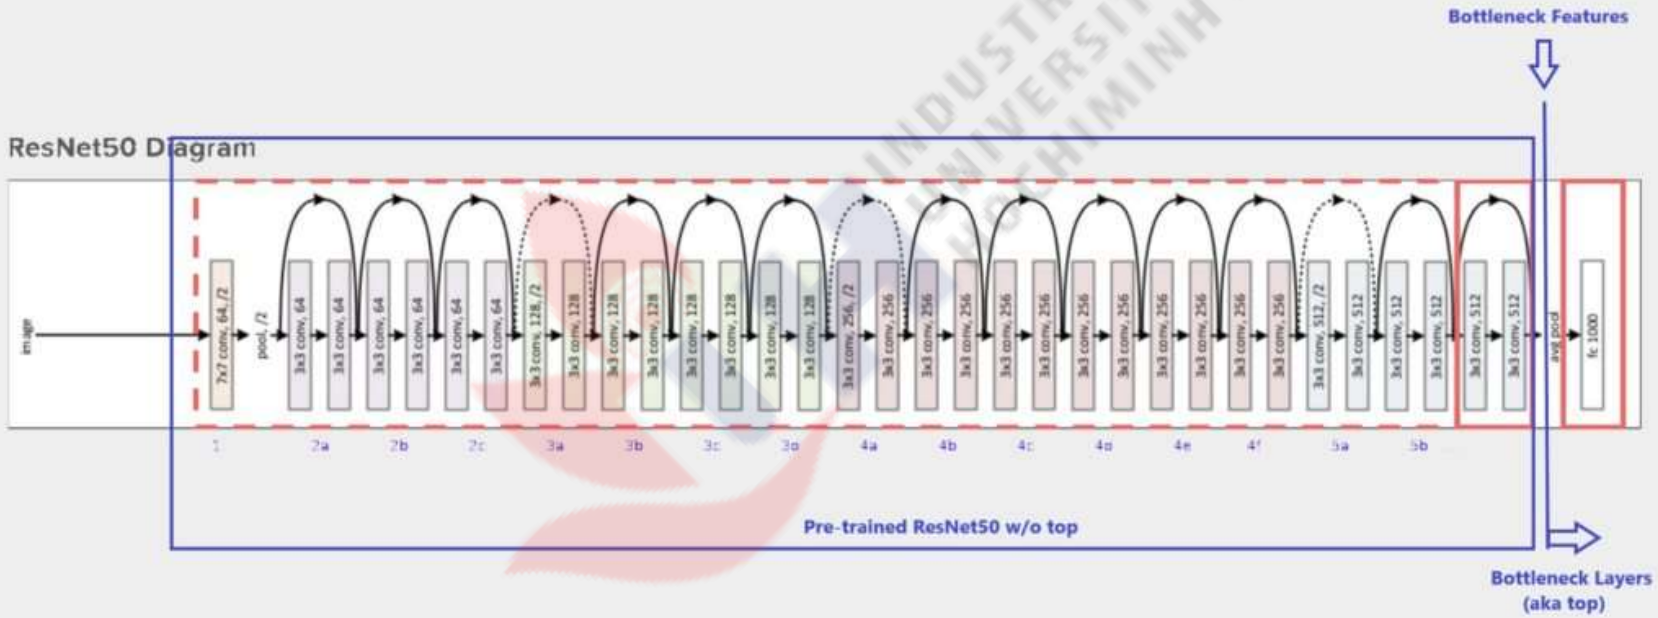

Thiết kế kiến trúc CNN

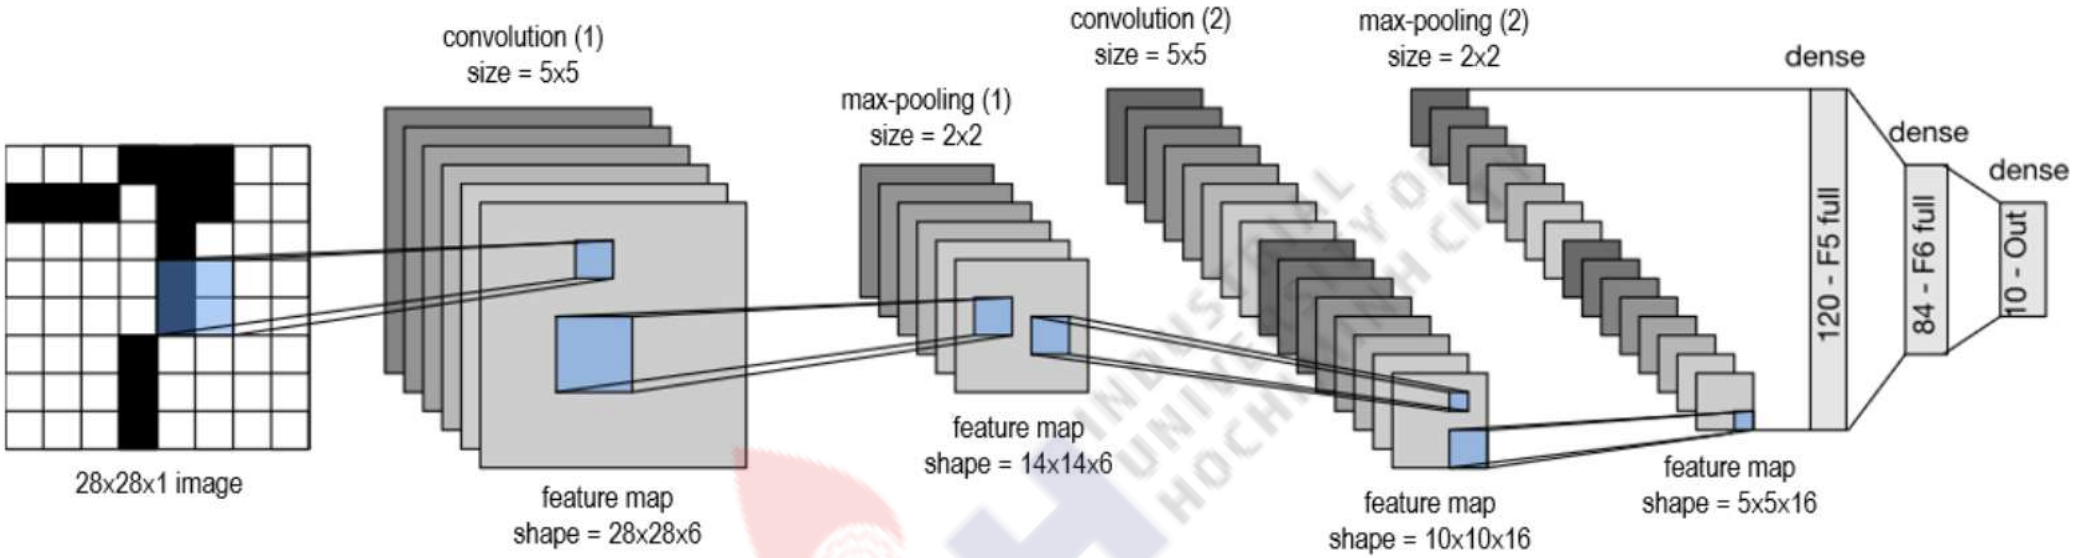

Thiết kế kiến trúc CNN

```
# model architecture
model = tf.keras.Sequential()

# input shape (28,28,1), convolution 1
model.add(tf.keras.layers.Conv2D(filters=6, kernel_size=3, padding='same',
activation='relu', input_shape=(28,28,1)))

# max pooling 1
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

# convolution 2
model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=5, activation='relu'))

# max pooling 2
model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

# Flatten
model.add(tf.keras.layers.Flatten())

# fully connected
model.add(tf.keras.layers.Dense(120, activation='relu'))
model.add(tf.keras.layers.Dense(84, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))

# Take a look at the model summary
model.summary()
```

```
Model: "sequential"
_________________________________________________________________
Layer (type)                 Output Shape              Param #
=================================================================
conv2d (Conv2D)              (None, 28, 28, 6)         156
_________________________________________________________________
max_pooling2d (MaxPooling2D) (None, 14, 14, 6)         0
_________________________________________________________________
conv2d_1 (Conv2D)            (None, 10, 10, 16)        2416
_________________________________________________________________
max_pooling2d_1 (MaxPooling2 (None, 5, 5, 16)          0
_________________________________________________________________
flatten (Flatten)            (None, 400)               0
_________________________________________________________________
dense (Dense)                (None, 120)               48120
_________________________________________________________________
dense_1 (Dense)              (None, 84)                10164
_________________________________________________________________
dense_2 (Dense)              (None, 10)                850
=================================================================
Total params: 61,706
Trainable params: 61,706
Non-trainable params: 0
_________________________________________________________________
```

Thiết kế kiến trúc CNN

- model.add(tf.keras.layers.Conv2D(filters=6, kernel_size=5, padding='same', activation='relu', input_shape=(28,28,1)))
- model.add(layers.Conv2D(6, (5, 5), activation='relu', input_shape=(28, 28, 1)))

=> Tạo 6 filter, mỗi filter kích thước 5x5. Tổng số weigh: 6x5x5= 150, bias weight : 6x1. Tổng cộng: 6x5x5+6x1=156.

padding='same' => không có padding, không thay đổi kích thước ma trận input = > output (28x28x6)

- model.add(tf.keras.layers.MaxPooling2D(pool_size=2))
- model.add(layers.MaxPooling2D((2, 2)))

=> Lớp maxPooling làm giảm kích thước ma trận theo 1 cửa sổ 2x2. Lớp này ko có weight. Từ input :(28x28x6) => output: (14x14x6)

- model.add(tf.keras.layers.Conv2D(filters=16, kernel_size=5, activation='relu'))
- model.add(layers.Conv2D(16, (5, 5), activation='relu'))

=>tạo 16 filter với kích thước mỗi filter là 5 x 5. Input : (14 x 14 x 6). Mỗi filter có 5x5 weights, lớp này có 5x5x6+1 = 151 weights cho mỗi phép tích chập. Tổng cộng có 16 filter: 151x16= 2416 weights.

padding = "SAME" => không có padding, output: (10x10x16)

- model.add(tf.keras.layers.MaxPooling2D(pool_size=2))

=> lớp maxPooling (2x2): output: (5x5x16), không có weight

- model.add(tf.keras.layers.Flatten())

=> sẽ flatten (5x 5 x 16) thành layer có 400 nodes, lớp này kết nối với FC có 120 nodes:

- model.add(tf.keras.layers.Dense(120, activation='relu'))

= 400x120+120=48120 weights

model.add(tf.keras.layers.Dense(84, activation='relu'))

= 120x84+84=10164

model.add(tf.keras.layers.Dense(84, activation='relu'))

=84x10+10=850

Tổng cộng: 61,706

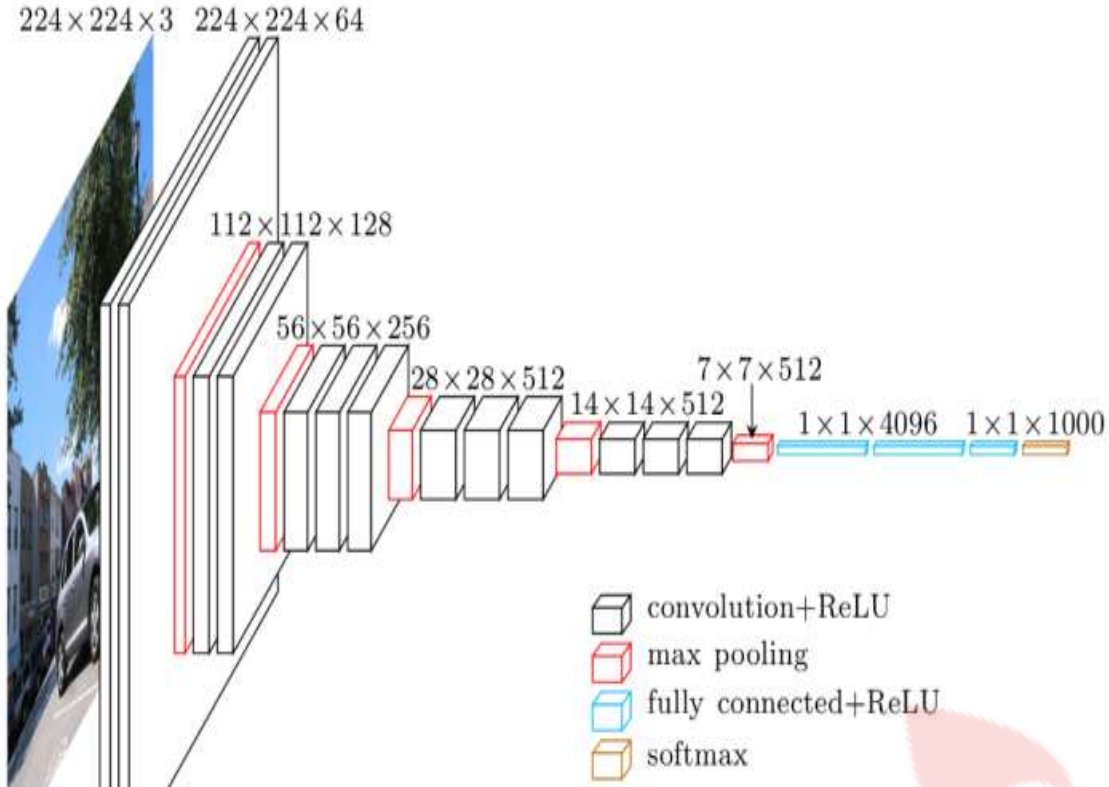


```
input_shape = (224, 224, 3)

#model
model = Sequential([
                Conv2D(64, (3, 3), input_shape=input_shape, padding='same',
activation='relu'),
                Conv2D(64, (3, 3), activation='relu', padding='same'),
                MaxPooling2D(pool_size=(2, 2), strides=(2, 2)),
                Conv2D(128, (3, 3), activation='relu', padding='same'),
                Conv2D(128, (3, 3), activation='relu', padding='same',),
                MaxPooling2D(pool_size=(2, 2), strides=(2, 2)),
                Conv2D(256, (3, 3), activation='relu', padding='same',),
                Conv2D(256, (3, 3), activation='relu', padding='same',),
                Conv2D(256, (3, 3), activation='relu', padding='same',),
                MaxPooling2D(pool_size=(2, 2), strides=(2, 2)),
                Conv2D(512, (3, 3), activation='relu', padding='same',),
                Conv2D(512, (3, 3), activation='relu', padding='same',),
                Conv2D(512, (3, 3), activation='relu', padding='same',),
                MaxPooling2D(pool_size=(2, 2), strides=(2, 2)),
                Conv2D(512, (3, 3), activation='relu', padding='same',),
                Conv2D(512, (3, 3), activation='relu', padding='same',),
                Conv2D(512, (3, 3), activation='relu', padding='same',),
                MaxPooling2D(pool_size=(2, 2), strides=(2, 2)),
                Flatten(),
                Dense(4096, activation='relu'),
                Dense(4096, activation='relu'),
                Dense(1000, activation='softmax')
])

model.summary()
```